AMZN Earnings Call – FY2025 Q1

Generated by discountingcashflows.com

Date: May 1, 2025

Operator

Thank you for standing by. Good day, everyone, and welcome to the Amazon.com First Quarter 2025 Financial Results Teleconference. At this time, all participants are in a listen-only mode. After the presentation, we will conduct a question-and-answer session. Today's call is being recorded. And for opening remarks, I will be turning the call over to the Vice President of Investor Relations, Mr. Dave Fildes. Thank you, sir. Please go ahead.

Dave Fildes

Hello and welcome to our Q1 2025 financial results conference call. Joining us today to answer your questions is Andy Jassy, our CEO; and Brian Olsavsky, our CFO. As you listen to today's conference call, we encourage you to have our press release in front of you, which includes our financial results, as well as metrics and commentary on the quarter. Please note, unless otherwise stated, all comparisons in this call will be against our results for the comparable period of 2024. Our comments and responses to your questions reflect management's views as of today, May 1, 2025 only, and will include forward-looking statements. Actual results may differ materially. Additional information about factors that could potentially impact our financial results is included in today's press release and our filings with the SEC, including our most recent annual report on Form 10-K and subsequent filings. During this call, we may discuss certain non-GAAP financial measures. In our press release, slides accompanying this webcast and our filings with the SEC, each of which is posted on our IR website, you will find additional disclosures regarding these non-GAAP measures, including reconciliations of these measures with comparable GAAP measures. Our guidance incorporates the order trends that we've seen to date, and what we believe today to be appropriate assumptions. Our results are inherently unpredictable and may be materially affected by many factors, including fluctuations in foreign exchange rates, changes in global economic and geopolitical conditions, tariff and trade policies, and customer demand and spending, including the impact of recessionary fears, inflation, interest rates, regional labor market constraints, world events, the rate of growth of the Internet, online commerce, cloud services and new and emerging technologies, and the various factors detailed in our filings with the SEC. Our guidance assumes, among other things. So, we don't conclude any additional business acquisitions, restructurings, or legal settlements. It's not possible to accurately predict demand for our goods and services, and therefore, our actual results could differ materially from our guidance. And now I'll turn the call over to Andy.

Andy Jassy

Thanks, Dave. Today we're reporting $165.7 billion in revenue, up 10% year-over-year, excluding the impact from foreign exchange rates. Operating income is $18.4 billion up 20% year-over-year and trailing 12 month free cash flow is $25.9 billion. We're pleased with our continued business progress, but more importantly with our pace of innovation and additional improvement in our customer experiences. In our stores business, we once again saw strong consumer resonance in our continued work on selection value and shipping speed. Our broad selection offers customers choice across their shopping journeys. We welcome well-known brands such as Oura Rings, Michael Kors, and The Ordinary, as well as a new shopping experience with Saks that offers a refined luxury assortment of fashion and beauty items from brands like Dolce&Gabbana, Balmain, Erdem, Giambattista Valli, and Jason Wu Collection. As always, we're working to keep prices low. And with this being an uncertain moment for consumers, it's even more important than it typically is. In Q1, we held deal events worldwide to help customers save over $500 million across the Big Spring Sale in the U.S. and Canada, Spring Deal Days in Europe, and Ramadan/Eid Sale events in Egypt, Saudi Arabia, Turkey, and UAE. Prime members will have more opportunities to save throughout the year including at our eleventh Prime Day event in July. Over the past few years, we've made significant progress in making our fulfillment network more efficient and cost effective. We've shared many times that an important turning point was regionalizing our national fulfillment network into regional hubs. By stocking items closer to where customers live, we're able to deliver more orders faster often in fewer packages and at lower delivery costs. The next challenge was getting as many items as possible into these regional nodes. Our inbound network which is how we get items to each fulfillment center hadn't been architected to leverage this new regionalization structure. So, we redesigned it and just rolled out a new inbound architecture that expands the share of products that we can place in each fulfillment center improving delivery speeds and lowering our cost to serve. In the first quarter, we once again set new delivery speed records with our fastest delivery ever for Prime members around the world and we delivered more items in the same day or next day in Q1 than any other quarter in our history. Looking ahead, we'll continue to refine our newly redesigned inbound network, build out our same day delivery sites and add additional robotics and automation throughout our buildings. You'll also see us expand the number of delivery stations that we have in rural areas of the U.S. So, we can get items to people who live in less densely populated areas much more quickly. I thought I'd share a few thoughts on the prospect of heightened tariffs on our stores business. Obviously none of us knows exactly where tariffs will settle or when. We haven't seen any attenuation of demand yet. To some extent, we've seen some heightened buying in certain categories that may indicate stocking up in advance of any potential tariff impact. We also have not seen the average selling price of retail items appreciably go up yet. Some of this reflects some forward buying we did in our first-party selling and some of that reflects some advanced inbounding our third-party sellers have done, but a fair amount of this is that most sellers just haven't changed pricing yet. Again, this could change depending on where tariffs settle. Amazon is not uniquely susceptible to tariffs. As it relates to China, retailers who aren't buying directly from China are typically buying from companies who themselves are buying from China, marking these items up, rebranding and selling to U.S. Consumers. These retailers are buying the product at a higher price than Chinese sellers selling directly to U.S. consumers in our marketplace. So, the total tariff will be higher for these retailers than for China Direct Sellers. It's also sometimes easy to forget what Amazon sells. We're not mostly selling high average selling price items, but we certainly sell a bunch. In the first quarter, our Everyday Essentials grew more than twice as fast as the rest of our business and represented one out of every three units sold in the U.S. and Amazon. Even if you exclude Whole Foods Market and Amazon Fresh, Amazon is one of the largest grocers in the U.S. with over $100 billion in gross sales last year. People are buying a lot of their everyday essentials at Amazon. We also have extremely large selection, hundreds of millions of unique SKUs, which means we're often able to weather challenging conditions better than others. When there are periods of discontinuity, substantial unexpected product trends emerge. Think about the pandemic, when items like masks and hand sanitizer became big sellers. When you have the broadest selection like we do and 2 million plus global sellers like we do, you're better positioned to help customers find whatever items matter to them at lower price points than elsewhere. Finally, when there are uncertain environments, customers tend to choose the provider they trust most. Given our really broad selection, low pricing and speedy delivery, we have emerged from these uncertain areas with more relative market segment share than we started and better set up for the future. I'm optimistic this could happen again. Moving to a few words on Amazon Ads, we're working hard to be the best place for brands of all sizes to grow their business. We are pleased with the strong growth on a very large base generating $13.9 billion of revenue in the quarter and growing 19% year-over-year. We're seeing strength across our broad portfolio of full funnel advertising offerings that help advertisers reach an average ad supported audience of more than 275 million in the U.S. alone. This includes our top of funnel efforts to drive brand awareness to bottom of funnel offerings where we measure outcomes at the point of conversion. Amazon Ads provides brands with tools to reach targeted audiences in our own entertainment properties such as Prime Video, Twitch and IMDB, in live sports such as NFL, NBA and NASCAR, audio content such as Amazon Music and Wondery and of course in our store as well as many other external sites such as Pinterest and BuzzFeed. All of our audience and measurement capabilities work for the ads we deliver across premium third party publishers through Amazon DSP and our secure clean rooms provide advertisers the ability to analyze data, produce core marketing metrics and understand how their marketing performs across various channels. We continue to see a lot of opportunity to further expand our full funnel capabilities for brands. AWS grew 17% year-over-year in Q1 and now sits at a $117 billion annualized revenue run rate. We continue to help organizations of all sizes accelerate their move to the cloud, helping to modernize their infrastructure, lower costs and speed up innovation. We signed new AWS agreements with companies including Adobe, Uber, Nasdaq, Ericsson, Fujitsu, Cargill, Mitsubishi Electric Corporation, General Dynamics Information Technology, GE Vernova, Booz Allen Hamilton, NextEra Energy, Publicis Sapient, Elastic, Netsmart, and many others. It's useful to remember that more than 85% of the global IT spend is still on premises, so not in the Cloud yet. It seems pretty straightforward to me that this equation will flip in the next 10 to 20 years. Before this generation of AI, we thought AWS had the chance to ultimately be a multi $100 billion revenue run rate business. We now think it could be even larger. If you believe your mission is to make customers' lives easier and better every day and you believe that every customer experience will be reinvented with AI, you're going to invest very aggressively in AI and that's what we're doing. You can see that in the thousand plus AI applications we're building across Amazon. You can see that with our next generation of Alexa named Alexa+. You can see that in how we're using AI in our fulfillment network, robotics, shopping, Prime Video and advertising experiences. And you can see that in the building blocks AWS is constructing for external and internal builders to build their own AI solutions. We're not dabbling here. We're very intentionally giving builders the broadest possible capabilities at every level of the AI stack cost effectively to use AI expansively across their businesses. At the bottom layer for those building models, our new custom AI chip Trainium 2 is starting to lay in capacity in larger quantities with significant appeal and demand. While we offer customers the ability to do AI in multiple chip providers and will for as long as I can proceed, customers doing AI at any significant scale realize that it can get expensive quickly. So, the 30% to 40% better price performance that Trainium 2 offers versus other GPU based instances is compelling. For AI to be as successful as we believe it can be, the price of inference needs to come down significantly. We consider this part of our mission and responsibility to help make it so. At the middle layer for those wanting to leverage frontier models to build Generative AI apps, Amazon Bedrock is our fully managed service that offers a choice of high performing foundation models with the most compelling set of features that make it easy to build high quality Generative AI applications. We continue to iterate quickly on Bedrock adding Anthropic's Claude 3.7 Sonnet hybrid reasoning model, their most intelligent model to date and Meta's Llama 4 family of models. We were also the first cloud service provider to make DeepSeek R1 and Mistral AI's Pixtral Large generally available as a fully managed model. And of course, we offer our own Amazon Nova state-of-the-art foundation models in Bedrock with the latest premier model launching yesterday. They deliver frontier intelligence and industry leading price performance and we have thousands of customers already using them including Slack, Siemens, Sumo Logic, Coinbase, FanDuel, Glean, and Blue Origin. A few weeks ago, we released the Amazon Nova Sonic, the new speech-to-speech foundation model that enables developers to build voice-based AI applications that are highly accurate, expressive, and human-like. Nova Sonic has lower word error rates and higher win rates over other comparable models for speech interactions. The technology was also abuzz by the potential of agents. To-date virtually all the agentic use cases have been of the question-answer variety. Our intention is for agents to perform wide-ranging complex, multi-step tasks by organizing a trip for setting the right temperature or music ambience in your house for dinner guests, or handling complex IT tasks to increase business productivity. There haven't been action-oriented agents like this until Alexa+. It was a technology to build these agents a still quite primitive and accurate, and requires constant human supervision. We just released a research preview of Amazon Nova Act, a new AI model trained to perform actions within a web browser, and it enables developers to breakdown complex workflows into reliable atomic commands like search, or checkout, or answer questions about the stream. It also enables them to add more detailed instructions to these commands, where needed, like don't accept the insurance upsell. Nova Act aims to move the current state-of-the-art accuracy of multi-step agentic actions from 30% to 60%, to 90-plus percent with the right set of building blocks to build these action-oriented agents. And the very top of the stack of the applications, this past quarter Q, the most capable generative AI assistance for accelerating software development, and leveraging your own data, plus a lightening-fast new agentic coding experience within the command line interface that could execute complex workflows autonomously. Customers are loving this. We also made generally available Give That Dudo with Amazon Q, enabling AI agents to assist multi-step tasks such as new feature development, code-based upgrades, or Java 8 and 11, while also offering code review and unit testing, all within the same familiar Give That platform. Our AI business has a multibillion dollar annual revenue run rate, continues to grow triple-digit year-over-year percentages, and is still in its very early days. While there is good reason for the high optimism about AI, I conclude my AWS comments with the reminder that there's still so much on-premises infrastructure yet to be moved to the cloud. Infrastructure modernization is much less actually to talk about the AI. The fundamental to any company's technology and invention capabilities to develop their productivity is speed and cost structure. If a company is to realize the full potential of AI, they are going to need their infrastructure and data in the cloud. I want to briefly mention a few other items. As I've referenced a couple of times, in Q1, we introduced Alexa+. Our next-generation of Alexa personal assistant was meaningfully smarter and more capable than our prior self, can both answer virtually any question and to take actions, and is free with Prime or available to non-Prime customers for $99 a month. We are just starting to roll this out in the U.S., and we'll be expanding to additional countries later this year. People are really liking Alexa+ thus far. We are excited and honored to be part of a joint venture that will be creating next-generation of the esteemed James Bond film franchise. We recently named the producer Amy Pascal and David Heyman to produce the next James Bond movie. Additionally, just a couple of days ago, Project Kuiper reached a significant milestone by launching our first satellite into orbit, with more being launched soon, and we expect to begin offering service to customers later this year. I'm proud of what our teams around the world have delivered. We are excited about what we are inventing and working on, as we speak. And with that, I'll turn it over to Brian for our financial update.

Brian Olsavsky

Thanks, Andy. I'll begin with our top line financial results. Worldwide revenue was $155.7 billion, a 10% increase year-over-year, excluding the impact of foreign exchange. This equates to $1.4 billion headwind from foreign exchange year-over-year in the quarter. Worldwide operating increase was $18.4 billion, approximately $400 million above the high-end of our guidance range. These results include one-time charges that impact North America and international operating income that I'll discuss in a moment. First, let's start with the net sales results for these segments. In the North America segment, first quarter revenue was $92.9 billion, an increase of 8% year-over-year. Our International segment revenue was $33.5 billion, also an increase of 8% year-over-year, excluding the impact of foreign exchange. Worldwide paid units were 8% year-over-year. Our priority is to provide value to our customers across our businesses. In the first quarter, we held multiple deal events around the world, which drove strong customer engagement. We saw broad-based strengths across our key business inputs, including record delivery speeds for Prime members, enabled by improved inventory placement. Our vast selection gives customers choice across various price points, particularly in categories like grocery, which includes everyday essentials. These are the items that people purchase most frequently. We partner with millions of independent sellers from around the world. These selling partners are important contributors to our broad selection, and worldwide third-party selling unit mix was 61% in Q1, consistent with Q1 of last year. Shifting to profitability, North America segment operating income was $5.8 billion, and international segment operating income was $1 billion, with operating margins of 6.3% in North America and 3% internationally. As I mentioned earlier, during the quarter, we've recorded one-time charges related to some historical customer returns that have not yet been resolved and some costs to receive inventory that was pulled forward into Q1 ahead of anticipated tariffs. Without these charges, North America and international operating margins would have been approximately 90 basis points and 70 basis points higher, with operating margins of 7.2% for North America and 3.7% for international. We are pleased with how our teams continue to execute and deliver for customers. In Q1, our newly re-architected inbound network drove productivity in our fulfillment and transportation network, leading to better inventory placement and higher units per package, and as a result, lower delivery costs. Beyond Q1, we have a number of initiatives underway to continue improving our cost structure, such as fine-tuning our inbound network, building out our same-day delivery sites, expanding our rural delivery network, and adding robotics and automation to our facilities. Better inventory placement remains a top priority. Better placement drives more in-stock selection, reduces travel distances, and speeds up delivery. And having inventory in the right place at the right time increases the likelihood that multiple items can be combined in a package, which helps reduce packaging and cost. Although progress won't always be linear, we have a good plan to continue to drive improvement over time. Shifting to advertising, advertising remains an important contributor to profitability in the North America and International segments. Advertising revenues grew 19% year-over-year. We're pleased with the accelerating growth on an increasingly large base. We're seeing strong adoption across our full-funnel advertising offering, as brands appreciate our ability to connect them with customers. We'll also continue to invest in other long-term opportunities. These efforts have the potential to be important to customers and Amazon in the future, including Kuiper, where we had our first launch of our production-designed satellites earlier this week. And we'll be launching more satellites throughout the year. We're closely monitoring the macroeconomic environment, including the impact of tariffs. We're planning for various outcomes, and we've taken a number of actions to protect the customer experience. We're doing everything we can to keep our prices low for customers in a way that makes economic sense. Moving next to our AWS segment, revenue was $29.3 billion in Q1, an increase of 17% year-over-year. AWS now has an annualized revenue run rate of more than $117 billion. During the first quarter, we continue to see growth in both Generative AI business and non-Generative AI offerings. As companies turn their attention to newer initiatives, bring more workloads to the cloud, restart or accelerate existing migrations from on-premises to the cloud, and tap into the power of Generative AI. AWS operating income was $11.5 billion and reflects our continued growth, coupled with our focus on driving efficiencies across the business. As we've said before, we expect AWS operating margins to fluctuate over time, driven in part by the level of investments we're making at any point in time. We plan to bring on an increasing amount of capacity in the back half of the year. Now, turning to our cash CapEx, which was $24.3 billion in Q1. The majority of the spend is to support the growing need for technology infrastructure. This primarily relates to AWS as we invest to support demand for our AI services, and increasingly in custom silicon like Trainium, as well as tech infrastructure to support our North America and international segments. We're also investing in our fulfillment and transportation network to support future growth and improve delivery speeds and our cost structure. This investment will support growth for many years to come. While we primarily focus our comments on operating income, I'd like to point out that our first quarter net income of $17.1 billion includes a pre-tax gain of $3.3 billion, included in non-operating income, and it relates to our investment in Anthropic. This activity is not related to Amazon's ongoing operations, but rather the result of the conversion of a portion of our convertible notes to non-voting preferred stock. Turning to Q2 guidance, as a reminder, our guidance considers a range of possibilities and takes into account Q1 results, trends in quarter day results, and expectations around the macroeconomic environment. Q2 net sales are expected to be between $159 billion and $164 billion. We estimate the year-over-year impact of changes in foreign exchange rates based on current rates, which we expect to be a headwind of approximately 10 basis points in the quarter. As a reminder, global currencies can fluctuate during the quarter. Q2 operating income is expected to be between $13 billion and $17.5 billion. This estimate includes the impact of our seasonal step-up in stock-based compensation expense in Q2, driven by the timing of our annual compensation cycle. The external environment remains complex, and as we have done throughout our history, we are focused on the inputs that we can control to protect the customer experience. We will work hard to remain the place customers trust for sharp prices, broad selection, and convenience. We'll remain focused on driving a better customer experience, and we still believe putting customers first is the only reliable way to create lasting value for our shareholders. With that, let's move on to your questions.

Operator

Thank you. At this time, we will now open the call up for questions. [Operator Instructions] Please hold while we poll for questions. And the first question comes from the line of Ross Sandler with Barclays. Please proceed with your question.

Ross Sandler

Great. I think I'm going to leave the China questions to others, and focus on AWS and kind of AI. So, Andy, it seems like you've been bringing on a lot more P5 GPU instances since February from what it looks like to kind of support all these new AI workloads. So, how would you characterize in the first quarter and maybe here in the second quarter, the kind of supply/demand imbalance that you talked about before around AI workloads? And when do you think that AWS will be in a position to kind of capture enough AI revenue to drive acceleration? Is that something that could happen this year? Do you see that more like next year, given your capacity constraints? Thank you very much.

Andy Jassy

Thanks, Ross. I would say that we've been bringing on a lot of P5s, which is a form of NVIDIA chip instances, as well as landing more and more Trainium 2 instances as fast as we can. And I would tell you that our AI business right now is a multi-billion dollar annual run rate business that's growing triple-digit percentages year-over-year. And we, as fast as we actually put the capacity in, it's being consumed. So, I think we could be driving -- we could be helping more customers and driving more revenue for the business if we had more capacity. We have a lot more Trainium 2 instances, and the next generation of NVIDIA's instances landing in the coming months. I expect that there are other parts of the supply chain that are a little bit jammed up as well at motherboards and some other componentry. And some of that is just because there is so much demand right now, but I do believe that the supply chain issues and the capacity issues will continue to get better as the year proceeds.

Operator

And the next question comes from the line of Eric Sheridan with Goldman Sachs. Please proceed with your question.

Eric Sheridan

Thanks so much for taking the question. Maybe I could ask a two-parter. First, in terms of strategy Andy, how do you think about positioning the company for the medium term, given all the levels of uncertainty out there about how the global trade environment might shift in the coming months? What do you see as the key strategic priorities that will allow the company to sort of be able to capitalize one way or another, depending on various elements of outcome? And how do you prioritize those investments in the months ahead? And then with respect to the one-quarter forward operating income guide, is there anything in there from a cost side that we should be thinking about as purely aligned with those types of investments against the trade landscape that might not repeat either later into this year or next year? Thanks so much.

Andy Jassy

Thanks, Eric. It's hard to tell what's going to happen with tariffs right now. It's hard to tell where they're going to settle and when they're going to settle. And so, a lot of what we're thinking about short and medium term actually turns out to be what we think about long-term too, which is, how do we actually have the broadest possible selection for customers at the lowest possible prices? And there's maybe never been a more important time in recent memory than trying to keep prices low, which we're heads down, pretty maniacally focused on, and then get things to people quickly and take care of customers. And that is the heart of what we're doing. And you can see different initiatives that we've taken within those priorities. We've done some forward buys of inventory where we're the first-party seller. Our third-party sellers have pulled forward a number of items so that they have inventory here as well. And those are all -- we're encouraging that because we're trying to keep prices as low as possible for customers. I think also, when you have as broad selection as we have, and we have much broader selection than other retailers, it means that when you've got this continuity, like we may potentially have, you're better able to help customers find what they want, no matter what those trends are. And I mentioned in my opening comments about what happened in the pandemic, and you can bet there are going to be things that we don't anticipate that customers really value and want that are different. It could be as simple, by the way, as just favoring other brands that maybe people didn't know about before, but where they have a more favorable price equation for customers. And I think when you've also got -- another thing that people forget is that when you've got 2 million-plus sellers, they're not all going to take the same strategy if there ends up being higher tariffs. I mean, there are going to be plenty of sellers that decide to pass on those higher costs to end consumers, but they're going to -- we have a lot of sellers in lots of different countries, and not all of them are going to pursue the same tack. And so, I think when you've got larger diversity like we have, we have a better chance of some of those sellers deciding that they're going to capture share, and they're not going to pass on all or any of those tariffs to customers. And so, I think customers are going to have a better chance of finding variety on selection and on lower prices when they come here. And the last thing I would say is that we have been in a number of our businesses, but just I'd say over the last 6 years or so, we have been diversifying where we produce things over a long period of time. We had, I would say, a meaningfully higher concentration of where we produced components for AWS or devices in China than we have now, where we've diversified meaningfully over that time. And we just thought that was wise to do so. And so, those are some of the ways that we're trying to make sure that we are protecting our customers over the short, medium, and long-term. But it turns out that most of those are the things that we focus most on, which is really broad selection for ultimate choice, really low prices, and very fast delivery.

Brian Olsavsky

Eric, I'll take your question on the guidance, and especially on operating income. I think that was your question, and the cost that might be in Q2. The thing I'd point to again is what I mentioned earlier, the stock-based comp always steps up generally in Q2 versus Q1, and then resets a rate that carries through for the next four quarters. You can look at historic trends to get an idea of that. Secondly, we do have some additional Kuiper launch costs in Q2. You saw a launch happen this week. And a reminder that we expense those launch costs until the point of commercialization, which a plan is to have that be later in this year.

Operator

Thank you. And the next question comes from the line of Justin Post with Bank of America. Please proceed with your question.

Justin Post

Great. Thanks. I'll go back to AWS. I know in the past you said revenues can be lumpy. Can you explain why they might fluctuate up and down if it's beyond just capacity? And you see the competitors with some pretty good growth rates. How do you think about the difference there? Obviously, your dollar growth is very good, but how do you think about the difference there versus some of your competitors? Thank you.

Andy Jassy

Well, the first thing I would say is when we've historically said that revenue can be lumpy, we've been saying this well before what's happened with AI over the last couple of years. And the reason for that is the sales cycle, particularly for enterprises. It's true for startups. What you really want is you want to have the type of capabilities where startups want to primarily choose to run on top of your platform. And that's true. If you look in the startup space, the vast majority of successful startups over the last 10 to 15 years have run on top of AWS. And it's unpredictable when those startups are going to find product market fit and grow substantially. And it's hard for them to predict and even harder for us to predict. And the same thing goes on the enterprise side, but in a different way. When the sales cycle on the enterprise side is that you spend time trying to convince people that they should move from on-premises to the cloud, and then that you have the right solution for them. And then, you pick a set of projects that you get experience on. And sometimes they use systems integrators. Sometimes they use our own professional services. Sometimes they're doing it themselves. Then there's a next tranche migrations. And those migrations just take time. And some companies get through them really quickly and some companies take longer to get through them. And what happens a lot of times too, is that they get excited and enthusiastic about the cost advantages and the speed of innovation advantages they get moving to the cloud. And what was supposed to be a smaller next tranche turns into a much larger next tranche. And all of that has been true for a long time. It's very hard for us to predict because it really is contingent on what enterprises, how they want to sequence it and resource it. Then you throw in AI, which has its own very fast growth cycle, particularly in certain types of use cases. And those change. I mean, I'll give you just some examples. In the early days, of the earliest days, I should say, of AI, what you've seen the most amount of has been initiatives that get you productivity and cost avoidance. And we've seen that from so many of our AWS customers. And we're doing a lot of it ourselves inside of Amazon using AI. And then, you've also seen, I would say, large scale training with a lot of those are running on top of this as well, and as Anthropic is building their next few training models on top of our Trainium 2 chip on AWS. And then, you've seen a couple of really big chatbots. And then, what you've seen just in the last few months is really kind of the explosion of coding agents. And if you just look at the growth of these coding agents, the last few months, these are companies like Cursor or Vercel, both of them run significantly on AWS, but just look at the growth of that over the last few months. You just couldn't have predicted that sort of growth. And so, that's why it's lumpy. Sometimes you'll have very significant increases that you didn't predict and you couldn't forecast. And then, they'll grow at a good rate, but maybe not the same rate before the next big kind of explosive growth. And I would tell you that everything I just mentioned, what's interesting in AI is that we still haven't gotten to all the other customer experiences that are going to be reinvented and all the other agents that are going to be built, they're going to take the role of a lot of different functions today. And even though we have a lot of combined inference in those areas, I would say we're not even at the second strike of the first batter in the first inning. It is so early right now. And then, I would just say on how to think about relative growth rates, you always have to -- the year-over-year growth rate is really only a function of the percentage growth on the base with which you are operating from. And we just have a very meaningfully larger base on the technology infrastructure side than others. And so, when it's still think about 17% year-over-year growth on a $117 billion revenue run rate, it's still pretty significant growth. And as I said, I think we could be doing more, more if we had more capacity and I expect that the capacity to ease in the coming months.

Operator

Thank you. Our next question comes from the line of Doug Anmuth with J.P. Morgan. Please proceed with your question.

Doug Anmuth

Thanks for taking the questions. One for Brian, one for Andy; Brian, just maybe to follow-up on AWS, but more on the margin side, we've seen a lot of fluctuation over the last couple of years and now hitting almost 40%. Maybe you can just talk about what's driving the outperformance and then how we should think about normalized margins going forward? And then, Andy, your comments on Alexa about moving to more complex tasks, can you talk about that more and just with Alexa, the products have been around for a long time, they've had different use cases. How do you get users to shift their behavior more with Alexa? Thanks.

Brian Olsavsky

Thanks, Doug. I'll take your first question. We did a strong quarter in AWS, as you mentioned, the margin performance. I would attribute it to the strong growth that we're seeing coupled with the impact of some continued investment we're making in innovation and technology, give you some examples. So, we invest in software and process improvements and it ends up optimizing our server capacity, which helps our infrastructure cost. We've been developing more efficient network using our low cost custom networking gear. We're working to maximize the power usage in our existing data centers, which both lowers our costs and also reclaims power for other newer workloads. And we're also seeing the impact of advancing custom silicon like Graviton. It provides lower costs not only for us, but also for our customers, the better price performance for them. But you're right, margins are impacted by a lot of things, including our level of investment, competitive pricing, the mix of Generative AI services as they're ramping up will continue to evolve over the years to come. So, we do have a lot of investment in infrastructure going on and planned for the second-half of the year. So, that will -- we'll start to see the impact of that. But we're happy with the performance of the team with generating cost savings. And it's a big focus as well as expanding the services and features for customers.

Andy Jassy

On the Alexa question, we're really excited about Alexa+. And as I mentioned earlier, it is -- she's much more intelligent, much more capable, and able to take real action. And to date, most of the agents that have been out there have really just been able to answer questions, which when it came out was very remarkable and -- but it's I think the future of agents is not just being intelligent, but also being able to take action. And that's actually, it requires a great model, but it also requires the ability to sync that model and to align that model with being able to take the right action and execute and implement the right APIs or you can have, very suboptimal results. And so, we've worked hard on that in Alexa+. We've been -- we started rolling out over the last several weeks. So, it's with now over 100,000 users, with more rolling out in the coming months. And so, far, the response from our customers has been very, very positive. People are excited about it. I think that it does a lot more things than what Alexa did before. And we're very fortunate in that we have over a half billion devices out there in people's homes and offices and cars. So, we have a lot of distribution already, but there will be to some degree, there will be a little bit of rewiring for people on what they can do because you get used to patterns. I mean, even the simple thing of not having to speak Alexa speak anymore where we're all used to saying Alexa before we want every, action to happen. And what you find is you really only have to do it the first time, and then really the conversation is ongoing where you don't have to say Alexa anymore. And I've been lucky enough to have the alpha and the beta, that I've been playing with for several months. And it took me a little bit of time to realize I didn't have to keep saying Alexa. And it's very freeing when you don't have to do that. And then, I think it's just experience in trying things. So, you can do things like you have guests coming over on a Saturday night for dinner and you can say Alexa, please, open the shades, put the lights on in the driveway and on the porch, increase the temperature five degree and pick music that would be great for dinner that's mellow. And she just does it. And like, when you have those types of experiences, it makes you want to do more of it. When I was in New York, when we were announcing, I asked her what were the -- we did the event way downtown. I asked her, what was great Italian restaurants or pizza restaurants. She gave me a list and then she asked me if she wanted me to make a reservation. I said yes and she made the reservation, confirmed the time. Like that, when you get into those types of routines and you have those types of experience, they're very, very useful. It is really like having a great personal assistant, which most people in the world don't have. And so, I think that the more and more that people get used to it, they will realize what she can do. And we're not going to be standing still. We have a lot more functionality that we plan to add in the coming months too.

Operator

Thank you. Our next question comes from the line of Brian Nowak with Morgan Stanley. Please proceed with your question.

Brian Nowak

Thanks for taking my questions. Hey, Andy and Brian. So, I have one for each of you guys. So, Andy, you have a very complicated retail business with a lot of moving pieces to it. I imagine you have a lot of good data on what you expect demand to be over the course of the year and the holidays and things. As you kind of step back and analyze the business and the tariff uncertainty, can you just sort of walk us through the one or two key areas operationally that you're most focused on just to ensure that Prime Day, Thanksgiving and the holidays go as smoothly as they possibly can? And then, secondly, on Brian, just to kind of go back to Eric's earlier question, as we think about the 2Q EBIT guide, are there any onetime costs or sort of tariff related costs in there similar to that $1 billion that you called out in the first quarter? Thanks.

Andy Jassy

Thanks, Brian. On the retail question, I would say that the areas that maybe we're most focused on to make sure we have not just a great Prime Day, but Prime Day is just one event as you know, and so is peak. We're trying to be great all year long for customers. The obvious ones are making sure that we help our sellers however we can, because there's uncertainty for our sellers as well. So, we're trying to make sure we provide a great experience. We're trying to make sure that we have the right diversity of sellers and low prices for our customers. I think all that you have to also be very thoughtful around how much inventory you bring into your fulfillment nodes at any one time, because you can imagine scenarios where either on your own when you're the first party seller or lots of third parties want to get as much inventory in as early as possible trying to beat a deadline on what may happen. And if you end up with too much inventory in your fulfillment network, it really slows down, your productivity and your ability to get things out as quickly as you want for customers at the cost structure you want. So, being able to manage that thoughtfully, we've learned that over the years, and I think the team is doing a really good job of balancing that right now.

Brian Olsavsky

And Brian, on your question about Q2, I guess I'd just reiterate what I had said earlier. We have stock based comp step up, which I think you can see the normal pattern for that if you look at our history. We have additional Kuiper expenses. Specifically, you asked about tariffs. We do have tariffs that we'll be paying on retail purchases based on current tariffs. It's not large in Q2. It's we had done a lot of pre-buying of inventory in Q1, as I mentioned earlier. But just generally, I think with the uncertainty, we've added a bit to the range that we've given you. We generally have a wide range, but just the general uncertainty that we're seeing and uncertainty of consumer demand and all the everything else is causing us to increase the ranges a bit. So, we'll see. We feel an informed view of Q2 right now. As Andy mentioned earlier, we saw, actually some strength in April based on what could end up being some pre buys of a number of things, but and advertising has been strong. So, we think there's a lot of positive trends, but certainly uncertainty right now for the quarter.

Operator

Thank you. And our final question will come from the line of Brent Thill with Jefferies. Please proceed with your question.

Brent Thill

Thanks. On AWS, I'm curious if you could give us the backlog number? And Andy, to your point about many of these core workload still yet to come to cloud, can you just update us on what you're seeing, are you seeing enterprise spring back, or are you seeing some confusion with AI clouding that transition and the timing of that? Just give some perspective on what you're seeing on migration. Thank you.

Dave Fildes

This is Dave. I'll just jump in on the backlog, and turn over to Andy. The backlog is $189 billion for Q1. That's up about 20% year-over-year. And the weighted average, meaning life on that is 4.1 years. And the AWS question around what we are seeing on the workloads that haven't moved, what I would say -- the way I would characterize it is that we were on a very aggressive March that was esthetical, but almost metronomic before the pandemic of enterprises deciding that they wanted to move off their on-premises infrastructure, because of the speed of innovation, the developer productivity, and the cost advantages of the cloud. And then, when you got into the pandemic, and the economy was concerned in the second-half of that few years, we had everybody trying to cost-optimize, including us by the way. And then, as we started to emerge from that trend, you saw generative AI explode, and everybody wanted to try to find a way to have a workload, or set a workload there, because people saw the potential, and it was also something that was generating a lot of interest publicly. And so, what we are seeing now over the last, call it, 16 to 18 months or so, is that enterprises realize they need to do both. And they want to do AI. Also to pilot at this point, AI, many of which will be successful, others will not be successful. The ones that are successful were scaled. But they also have a lot more initiatives that they still haven't gotten to on the AI side, either because the willingness skill set, where they're picking the first set to get experience with, or as they're waiting to see the cost inference continue to go down, which it will. I mean, you will not get the expensiveness that we all know is coming in AI until we keep getting the cost inference down, even though it's going like crazy right now. But at the same time, I would say that we see an increased resurgence, and understanding for enterprises that they're dropping the low-hanging fruit if they don't move their infrastructure to the cloud, for all the reasons I mentioned earlier. So, you're starting to see those plans pick up again. As I mentioned earlier to one of the questions, you know the side that you're going to transform your infrastructure from on-premises and cloud, and see it happen in there months. That is typically a multiyear process that some companies do it fast, some companies do it slower, but they do it cautious and they do it softly because they can't afford for their application not to work if they're making that transition. And we are having meaningful success in those conversations, and in companies choosing to transform their infrastructure on top of AWS, and I think you'll see that moving forward too. Thank you for your time joining us today, and for your questions. A replay will be available on the Investor Relations website for at least three months. We appreciate your interest in Amazon, and we look forward to talking with you again next quarter.

Operator

And ladies and gentlemen, that does conclude today's teleconference. You may disconnect your lines at this time. Thank you for your participation.