AVGO Earnings Call – FY2024 Q2

Date: June 12, 2024

Operator

Welcome to Broadcom Inc. Second Quarter Fiscal Year 2024 Financial Results Conference Call. At this time, for opening remarks and introductions, I would like to turn the call over to Ji Yoo, Head of Investor Relations of Broadcom Inc.

Ji Yoo

Thank you, Operator, and good afternoon, everyone. Joining me on today's call are Hock Tan, President and CEO; Kirsten Spears, Chief Financial Officer and Charlie Kawwas, President, Semiconductor Solutions Group. Broadcom distributed a press release and financial tables after the market closed, describing our financial performance for the second quarter of fiscal year 2024. If you did not receive a copy, you may obtain the information from the Investor section of Broadcom's website at Broadcom.com. This conference call is being webcast live and an audio replay of the call can be accessed for 1 year through the Investor section of Broadcom's website. During the prepared comments, Hock and Kirsten will be providing details of our second quarter fiscal year 2024 results, guidance for our fiscal year 2024, as well as commentary regarding the business environment. We'll take questions after the end of our prepared comments. Please refer to our press release today and our recent filings with the SEC for information on the specific risk factors that could cause our actual results to differ materially from the forward-looking statements made on this call. In addition to US GAAP reporting, Broadcom reports certain financial measures on a non-GAAP basis. A reconciliation between GAAP and non-GAAP measures is included in the tables attached to today's press release. Comments made during today's call will primarily refer to our non-GAAP financial results. I'll now turn the call over to Hock.

Hock Tan

Thank you, Ji. And thank you everyone for joining today. In our fiscal Q2 2024 results -- consolidated net revenue was $12.5 billion, up 43% year-on-year as revenue included a full quarter of contribution from VMware. But if we exclude VMware, consolidated revenue was up 12% year-on-year. And this 12% organic growth in revenue was largely driven by AI revenue, which stepped up 280% year-on-year to $3.1 billion, more than offsetting continued cyclical weakness in semiconductor revenue from enterprises and telcos. Let me now give you more color on our two reporting segments. Beginning with software. In Q2 infrastructure software segment revenue of $5.3 billion was up 175% year-on-year and included $2.7 billion in revenue contribution from VMware, up from $2.1 billion in the prior quarter. The integration of VMware is going very well. Since we acquired VMware, we have modernized the product SKUs from over 8,000 disparate SKUs to four core product offerings and simplified the go-to-market flow, eliminating a huge amount of channel conflicts. We are making good progress in transitioning all VMware products to a subscription licensing model. And since closing the deal, we have actually signed up close to 3,000 of our largest 10,000 customers to enable them to build a self-service virtual private cloud on-prem. Each of these customers typically sign up to a multi-year contract, which we normalize into an annual measure known as Annualized Booking Value, or ABV. This metric, ABV for VMware products, accelerated from $1.2 billion in Q1 to $1.9 billion in Q2. By reference, for the consolidated Broadcom software portfolio, ABV grew from $1.9 billion in Q1 to $2.8 billion over the same period in Q2. Meanwhile, we have integrated SG&A across the entire platform and eliminated redundant functions. Year-to-date, we have incurred about $2 billion of restructuring and integration costs and drove our spending run rate at VMware to $1.6 billion this quarter, from what used to be $2.3 billion per quarter pre-acquisition. We expect spending will continue to decline towards a $1.3 billion run rate exiting Q4, better than our previous $1.4 billion plan, and will likely stabilize at $1.2 billion post-integration. VMware revenue in Q1 was $2.1 billion, grew to $2.7 billion in Q2, and will accelerate towards a $4 billion per quarter run rate. We therefore expect operating margins for VMware to begin to converge towards that of classic Broadcom software by fiscal 2025. Turning to semiconductors, Let me give you more color by end markets. Networking. Q2 revenue of $3.8 billion grew 44% year-on-year, representing 53% of semiconductor revenue. This was again driven by strong demand from hyperscalers for both AI networking and custom accelerators. It's interesting to note that as AI data center clusters continue to deploy, our revenue mix has been shifting towards an increasing proportion of networking. We doubled the number of switches we sold year-on-year, particularly the PAM-5 and Jericho3, which we deployed successfully in close collaboration with partners like Arista Networks, Dell, Juniper, and Supermicro. Additionally, we also double our shipments of PCI Express switches and NICs in the AI backend fabric. We're leading the rapid transition of optical interconnects in AI data centers to 800 gigabit bandwidth, which is driving accelerated growth for our DSPs, optical lasers, and PIN diodes. And we are not standing still. Together with these same partners, we are developing the next generation switches, DSP, and optics that will drive the ecosystem towards 1.6 terabit connectivity to scale out larger AI accelerated clusters.

,

Next year, we expect all mega-scale GPU deployments to be on Ethernet. We expect the strength in AI to continue, and because of that, we now expect networking revenue to grow 40% year-on-year compared to our prior guidance of over 35% growth. Moving to wireless. Q2 wireless revenue of $1.6 billion grew 2% year-on-year, was seasonally down 19% quarter-on-quarter and represents 22% of semiconductor revenue. And in fiscal '24, helped by content increases, we reiterate our previous guidance for wireless revenue to be essentially flat year-on-year. This trend is wholly consistent with our continued engagement with our North American customer, which is deep, strategic, and multiyear and represents all of our wireless business. Next, our Q2 server storage connectivity revenue was $824 million or 11% of semiconductor revenue, down 27% year-on-year. We believe though, Q2 was the bottom in server storage. And based on updated demand forecast and bookings, we expect a modest recovery in the second half of the year. And accordingly, we forecast fiscal '24 server storage revenue to decline around the 20% range year-on-year. Moving on to broadband. Q2 revenue declined 39% year-on-year to $730 million and represented 10% of semiconductor revenue. Broadband remains weak on the continued pause in telco and service provider spending. We expect Broadcom to bottom in the second half of the year with a recovery in 2025. Accordingly, we are revising our outlook for fiscal '24 broadband revenue to be down high 30s year-on-year from our prior guidance for a decline of just over 30% year-on-year. Finally, Q2 industrial rev -- resale of $234 million declined 10% year-on-year. And for fiscal '24, we now expect industrial resale to be down double-digit percentage year-on-year compared to our prior guidance for high single-digit decline. So to sum it all up, here's what we are seeing. For fiscal '24, we expect revenue from AI to be much stronger at over $11 billion. Non-AI semiconductor revenue has bottomed in Q2 and is likely to recover modestly for the second half of fiscal '24. On infrastructure software, we're making very strong progress in integrating VMware and accelerating its growth. Pulling all these three key factors together, we are raising our fiscal '24 revenue guidance to $51 billion. And with that, let me turn the call over to Kirsten.

Kirsten Spears

Thank you, Hock. Let me now provide additional detail on our Q2 financial performance, which included a full quarter of contribution from VMware. Consolidated revenue was $12.5 billion for the quarter, up 43% from a year ago. Excluding the contribution from VMware, Q2 revenue increased 12% year-on-year. Gross margins were 76.2% of revenue in the quarter. Operating expenses were $2.4 billion and R&D was $1.5 billion, both up year-on-year primarily due to the consolidation of VMware. Q2 operating income was $7.1 billion and was up 32% from a year ago with operating margin at 57% of revenue. Excluding transition costs, operating profit of $7.4 billion was up 36% from a year ago, with operating margin of 59% of revenue. Adjusted EBITDA was $7.4 billion or 60% of revenue. This figure excludes $149 million of depreciation. Now a review of the P&L for our two segments, starting with semiconductors. Revenue for our semiconductor solutions segment was $7.2 billion and represented 58% of total revenue in the quarter. This was up 6% year-on-year. Gross margins for our semiconductor solutions segment were approximately 67%, down 370 basis points year-on-year, driven primarily by a higher mix of custom AI accelerators. Operating expenses increased 4% year-on-year to $868 million on increased investment in R&D, resulting in semiconductor operating margins of 55%. Now moving on to infrastructure software. Revenue for infrastructure software was $5.3 billion, up 170% year-on-year, primarily due to the contribution of VMware and represented 42% of revenue. Gross margin for infrastructure software were 88% in the quarter, and operating expenses were $1.5 billion in the quarter, resulting in infrastructure software operating margin of 60%. Excluding transition costs, operating margin was 64%. Now moving on to cash flow. Free cash flow in the quarter was $4.4 billion and represented 36% of revenues. Excluding cash used for restructuring and integration of $830 million, free cash flows of $5.3 billion were up 18% year-on-year and represented 42% of revenue. Free cash flow as a percentage of revenue has declined from 2023 due to higher cash interest expense from debt related to the VMware acquisition and higher cash taxes due to a higher mix of US income and the delay in the reenactment of Section 174. We spent $132 million on capital expenditures. Days sales outstanding were 40 days in the second quarter, consistent with 41 days in the first quarter. We ended the second quarter with inventory of $1.8 billion down 4% sequentially. We continue to remain disciplined on how we manage inventory across our ecosystem. We ended the second quarter with $9.8 billion of cash and $74 billion of gross debt. The weighted average coupon rate and years to maturity of our $48 billion in fixed rate debt is 3.5% and 8.2 years respectively. The weighted average coupon rate and years to maturity of our $28 billion in floating rate debt is 6.6% and 2.8 years, respectively. During the quarter, we repaid $2 billion of our floating rate debt, and we intend to maintain this quarterly repayment of debt throughout fiscal 2024. Turning to capital allocation. In the quarter, we paid stockholders $2.4 billion of cash dividends based on a quarterly common stock cash dividend of $5.25 per share. In Q2, non-GAAP diluted share count was 492 million as the 54 million shares issued for the VMware acquisition were fully weighted in the second quarter. We paid $1.5 billion withholding taxes due on vesting of employee equity, resulting in the elimination of 1.2 million AVGO shares. Today, we are announcing a 10-for-1 forward stock split of Broadcom's common stock to make ownership of Broadcom stock more accessible to investors and to employees. Our stockholders of record after the close of market on July 11, 2024, will receive an additional nine shares of common stock after the close of market on July 12, with trading on a split-adjusted basis expected to commence at market open on July 15, 2024. In Q3, reflecting a post-split basis, we expect share count to be approximately 4.92 billion shares. Now on to guidance. We are raising our guidance for fiscal year 2024 consolidated revenue to $51 billion and adjusted EBITDA to 61%. For modeling purposes, please keep in mind that GAAP net income and cash flows in fiscal year 2024 are impacted by restructuring and integration-related cash costs due to the VMware acquisition. That concludes my prepared remarks. Operator, please open up the call for questions.

Operator

Thank you. [Operator Instructions] And our first question will come from the line of Vivek Arya with Bank of America. Your line is open.

Vivek Arya

Thanks for taking my question. Hock, I would appreciate your perspective on the emerging competition between Broadcom and NVIDIA across both Accelerators and Ethernet switching. So on the Accelerator side, they are going to launch their Blackwell product that many of the same customers that you have a very large position in the custom compute. So I'm curious how you think customers are going to do that allocation decision, just broadly what the visibility is. And then I think Part B of that is as they launch their Spectrum-X Ethernet switch, do you think that poses increasing competition for Broadcom and the Ethernet switching side in AI for next year? Thank you.

Hock Tan

Very interesting question, Vivek. On AI accelerators, I think we are operating on a different -- to start with scale, much as a different model. It is -- and on the GPUs, which are the AI accelerator of choice on merchant -- in a merchant environment is something that is extremely powerful as a model. It's something that NVIDIA operates in, in a very, very effective manner. We don't even think about competing against them in that space, not in the least. That's where they're very good at and we know where we stand with respect to that. Now what we do for very selected or selective hyperscalers is, if there's a scale and the skills to try to create silicon solutions, which are AI accelerators to do particular very complex AI workloads. We are happy to use our IP portfolio to create those custom ASIC AI accelerator. So I do not see them as truly competing against each other. And far for me to say I'm trying to position myself to be a competitor on basically GPUs in this market. We're not. We are not a competitor to them. We don't try to be, either. Now on networking, maybe that's different. But again people may be approaching and they may be approaching it from a different angle. We are as I indicated all along, very deep in Ethernet as we've been doing Ethernet for over 25 years, Ethernet networking. And we've gone through a lot of market transitions, and we have captured a lot of market transitions from cloud-scale networking to routing and now AI. So it is a natural extension for us to go into AI. We also recognize that being the AI compute engine of choice in merchants in the ecosystem, which is GPUs, that they are trying to create a platform that is probably end-to-end very integrated. We take the approach that we don't do those GPUs, but we enable the GPUs to work very well. So if anything else, we supplement and hopefully complement those GPUs with customers who are building bigger and bigger GPU clusters.

Vivek Arya

Thank you.

Operator

One moment for our next question, and that will come from the line of Ross Seymore with Deutsche Bank. Your line is open.

Ross Seymore

Hi guys. Thanks for taking my question. I want to stick on the AI theme, Hock. The strong growth that you had in the quarter, the 280% year-over-year, could you delineate a little bit between if that's the compute offload side versus the connectivity side? And then as you think about the growth for the full year, how are those

split in that realm as well? Are they kind of going hand-in-hand? Or is one side growing significantly faster than the other, especially with the I guess, you said the next-generation accelerators are now going to be Broadcom as well?

Hock Tan

Well, to answer your question on the mix, you are right. It's something we don't really predict very well, not understand completely except in hindsight. Because it's like, to some extent, to the cadence of deployment of when they put in the AI accelerators versus when they put in the infrastructure that puts it together, the networking. And we don't really quite understand it 100%. All we know, it used to be 80% accelerators, 20% networking. It's now running closer to two-thirds accelerators, one-thirds networking and we'll probably head towards 60%-40% by the close of the year.

Ross Seymore

Thank you.

Operator

Thank you. One moment for our next question. And that will come from the line of Stacy Rasgon with Bernstein. Your line is open.

Stacy Rasgon

Hi, guys. Thanks for taking my question. I wanted to ask about the $11 billion AI guide. You'd be at $11.6 billion even if you didn't grow AI from the current level in the second half. And it feels to me like you're not suggesting that. It feels to me like you think you could be [guided] (ph). So why wouldn't that AI number be a lot more than $11.6 billion? It feels like it ought to be. Or am I missing something?

Hock Tan

Because I guided just over $11 billion, Stacy could be what you think it is. It's -- quarterly shipments get sometimes very lumpy. And it depends on rate of deployment, depends on a lot of things. So you may be right. You may estimate it better than I do, but the general trajectory is getting better.

Stacy Rasgon

Okay. So I guess again, how do I -- are you just suggesting that, that more than $11 billion is sort of like the worst it could be because that would just be flat at the current levels, but you're also suggesting that things are getting better into the back half?

Hock Tan

Correct.

Stacy Rasgon

Okay. So I guess we just take that, that's a very -- if I'm reading it wrong, that's just a very conservative number?

Hock Tan

That's the best forecast I have at this point, Stacy.

Stacy Rasgon

All right. Okay, Hock, thank you. I appreciate it.

Hock Tan

Thank you.

Operator

One moment for our next question, and that will come from the line of Harlan Sur with JPMorgan. Your line is open.

Harlan Sur

Yeah, good afternoon. Thanks for taking my question. Hock, on cloud and AI networking silicon, good to see that the networking mix is steadily increasing. Like clockwork, the Broadcom team has been driving a consistent two year cadence, right of new product introductions, Trident, Tomahawk, Jericho family of switching and routing products for the past seven generations. You layer on top of that your GPU -- TPU customers are accelerating their cadence of new product introductions and deployments of their products. So is this also driving faster adoption curve for your latest Tomahawk and Jericho products? And then maybe just as importantly, like clockwork, it is been two years since you've introduced Tomahawk 5 product introduction, right which if I look back historically, means you have silicon and are getting ready to introduce your next-generation three-nanometer Tomahawk 6 products, which would, I think, puts you two years to three years ahead of your competitors. Can you just give us an update there?

Hock Tan

Harlan, you're pretty insightful there. Yes, we launched Tomahawk 5 in 2023. So you're right, by late 2025, the time we should be coming out with Tomahawk 6, which is the 100 terabit switch, yes.

Harlan Sur

And is the -- is this acceleration of cadence by your GPU and TPU partners, is that also what's kind of driving the strong growth in the networking products?

Hock Tan

Well, you know what, sometimes you have to let things take its time. But it's two-year cadence so we're right on. Late 2023, once when we shoot it out to a Tomahawk 5 and adoption. You're correct with AI has been tremendous because it ties in with the need for a very large bandwidth in the networking, in the fabric for AI clusters, AI data centers. But regardless, we have always targeted Tomahawk 6 to be out two years after that, which should put it into late '25.

Harlan Sur

Okay, thank you Hock.

Operator

Thank you. One moment for our next question, and that will come from the line of Ben Reitzes with Melius. Your line is open.

Ben Reitzes

Hi, thanks a lot. And congrats on the quarter and guide. Hock, I wanted to talk a little bit more about VMware. Just wanted to clarify if it is indeed going better than expectations. And how would you characterize the customer willingness to move to subscription? And also just a little more color on Cloud Foundation. You've cut the price there, and are you seeing that beat expectations? Thanks a lot.

Hock Tan

Thanks, and thanks for your kind regards on the quarter. But it's -- as far as VMware is concerned, we're making good progress. The journey is not over by any means, but it is pretty much very much to expectation. Moving to subscription, well, in VMware we are very slow compared to, I mean a lot of other guys, Microsoft, Salesforce, Oracle, who have already been pretty much in subscription. So VMware is late in that process. But we're trying to make up for it by offering it and offering it in a very, very compelling manner because subscription is the right things to do, right? It's a situation where you put out your product offering, and you update it, patch it, but update it feature-wise, everything as capabilities on a continual basis, almost like getting your news on an ongoing basis, subscription online versus getting it in a printed manner once a week. That's how I compare perpetual to subscription. So it is very interesting for a lot of people to want to can't get on. And so to no surprise, they are getting on very well. The big selling point we have as I indicated, is the fact that we are not just trying to keep customers kind of stuck on just server or compute virtualization. That's a great product, great technology, but it's been out for 20 years. Based on what we are offering now at a very compelling price point, compelling in a very attractive price point, the whole stack, software stack to use vSphere and its basic fundamental technology to virtualize networking, storage, operation and management, the entire data center and create this self-service private cloud. And thanks for saying it, you're right, and we have priced it down to the point where it is comparable with just compute virtualization. So yes, that is getting a lot of interest, a lot of attention from the customers. We have signed up who would like to deploy -- the ability to deploy private cloud, their own private cloud on-prem. As a nice complement, maybe even alternative or hybrid to public clouds, that's the selling point, and we are getting a lot of interest from our customers in doing that.

Ben Reitzes

Great. And it's on track for $4 billion by the fourth quarter still, which is reiterated?

Hock Tan

Well, I didn't give a specific time frame, did I? But it's on track as we see this process growing towards a $4 billion quarter.

Ben Reitzes

Okay, thanks a lot Hock.

Hock Tan

Thanks.

Operator

Thank you. One moment for our next question, and that will come from the line of Toshiya Hari with Goldman Sachs. Your line is open.

Toshiya Hari

Hi, thank you so much for taking my question. I guess kind of a follow-up to the previous question on your software business. Hock, you seem to have pretty good visibility into hitting that $4 billion run rate over the medium term, perhaps. You also talked about your operating margin in that business converging to classic Broadcom levels. I know the integration is not done and you're still kind of in debt paydown mode. But how should we think about your growth strategy beyond VMware? Do you think you have enough drivers, both on the semiconductor side and the software side to continue to drive growth or is M&A still an option beyond VMware? Thank you.

Hock Tan

Interesting question. And you're right. As I indicated in my remarks, even we found the contribution from VMware this past quarter where we have AI helping us, but we have non-AI semiconductor sort of bottoming out. We're able to show 12% organic growth year-on-year. So almost have to say, so do we need to rush to buy another company? Answer is no. But all options are always open because we are trying to create the best value for our shareholders who have entrusted us with the capital to do that. So I would not discount that alternative because our strategy, our long-term model has always been to grow through a combination of

acquisition, but also on the assets we acquire to really improve, invest, and operate them better to show organic growth as well. But again, organic growth often enough is determined very much by how fast your market would grow. So we do look towards acquisitions now and then.

Toshiya Hari

All right. Thank you.

Operator

Thank you. One moment for our next question, and that will come from the line of Blayne Curtis with Jefferies. Your line is open.

Blayne Curtis

Hi, thanks for taking my question. I wanted to ask you Hock, on the networking business kind of ex AI. Obviously, I think there's an inventory correction the whole industry is seeing. But just kind of curious, I don't think you mentioned that it was at a bottom. So just the perspective, I think it's down about [60%] (ph) year-over-year. Is that business finding a bottom? I know you said overall whole semi business should -- non-AI should see a recovery. Are you expecting any there any perspective on just customer inventory levels in that segment?

Hock Tan

We see it behaving. I didn't particularly call it out, obviously because more than anything else, I kind of link it very much to server storage, non-AI that is. And we call server storage as at a bottom Q2, and we call it to recover modestly second half of the year. We see the same thing in networking, which is a combination of enterprise networking, as well as the hyperscalers who run their traditional workloads on those, though it's hard to figure out sometimes. But it is. So we see the same trajectory as we are calling out on server storage.

Blayne Curtis

Okay, thank you.

Operator

Thank you. One moment for our next question, and that will come from the line of Timothy Arcuri with UBS. Your line is open. Mr. Arcuri, your line is open.

Timothy Arcuri

Hi, sorry. Thanks. Hock, is there a way to sort of map GPU demand back to your AI networking opportunity? I think I've heard you say in the past that if you spent $10 billion on GPU compute, you need to spend another $10 billion on other [infrastructure] (ph), most of which is networking. So I'm just kind of wondering if when you see these big GPU numbers, is there sort of a rule of thumb that you use to map it back to what the opportunity will be for you? Thanks.

Hock Tan

There is, but it's so complex, I stopped creating such a model, Tim. I've said it. But there is because one would say that for every -- you almost say, for every $1 billion you spend on GPU, you probably would spend probably on networking, and if we include the optical interconnects as part of it, though we are not totally in that market, except for the components like DSPs, lasers, PIN diodes that go into those, high-bandwidth optical connect. But if you just take optical connects in totality, switching, all the networking components, it goes into -- attaches itself to clustering a bunch of GPUs, you probably would say that about 25% of the value of the GPU goes to networking, the rest on networking. Now not entirely all of it is my available market. I don't do the optical connects, but I do the few components I talked about in it. But roughly, the simple way to look at it is probably about 25%, maybe 30% of all these infrastructure components is kind of attached to the GPU value point itself. But having said that, it's never – we are never that precise that deployment is the same way. So you may see the deployment of GPU or purchase of GPU much earlier. And the networking comes later or sometimes less the other way around, which is why you're seeing the mix going on within my AI revenue mix. But typically, you run towards that range over time.

Timothy Arcuri

Perfect Hock, thank you so much.

Operator

Thank you. One moment for our next question, and that will come from the line of Thomas O'Malley with Barclays. Your line is open.

Thomas O'Malley

Hi, guys. Thanks for taking my question. And nice results. My question in regards to the custom ASIC AI. Hock, you had a long run here of a very successful business, particularly with one customer. If you look in the market today, you have a new entrant who's playing with different customers. And I know that you said historically, that's not really a direct customer to you. But could you talk about what differentiates you from the new entrant in the market as of late? And then there's been profitability questions around the sustainability of gross margins longer term. Can you talk about if you see any increased competition? And if there's really areas that you would deem more or less defensible in your profile today? And if you would see kind of that additional entrant maybe attack any of those in the future?

Hock Tan

Let me take the second part first, which is our AI -- custom AI accelerator business. It is a very profitable business, and let me put to scale -- look examine from a model point of view. I mean, each of these AI accelerators no different from a GPU. The way these large language models get run computing, get run on these accelerators, no one single accelerator, as you know, can run these big large language models. You need multiple of it no matter how powerful those accelerators

are. But also, and the way the models are run, there is a lot of memory access to memory requirements. So each of this accelerator comes with a large amount of cache memory, as you call it, what you guys probably now know as HBM, high-bandwidth memory specialized for AI accelerators or GPUs. So we're supplying both in our custom business. And the logic side of it, where the compute function is on doing the chips, the margin there are no different than the margin in any -- in most of any of a semiconductor silicon chip business. But when you're attached to it, a huge amount of memory, memory comes from a third-party. There are a few memory makers who make this specialized thing. We don't do margin stacking on that part. So by almost buying basic math will dilute the margin of these AI accelerators when you sell them with memory, which we do. It does push up revenue somewhat higher but it has diluted the margin. But regardless, the spend, the R&D, the OpEx that goes to support this as a percent of the revenue, which is higher revenue, so much less. So on an operating margin level, this is easily as profitable, if not more profitable, given the scale that each of those custom AI accelerators can go up to. It's even better than our normal operating margin scale. So that's the return on investment that attracts and keeps us going at this game. And this is more than a game. It is a very difficult business. And to answer your first question, there is only one Broadcom, period.

Thomas O'Malley

Thanks Hock.

Operator

Thank you. One moment for our next question, and that will come from the line of Karl Ackerman with BNP. Your line is open.

Karl Ackerman

Hi, thank you. Good afternoon. Hock, your networking switch portfolio with Tomahawk and Jericho chipsets allows hyperscalers to build AI clusters using either a switch-scheduled or endpoint-scheduled network. And that, of course is unique among competitors. But as hyperscalers seek to deploy their own unique AI clusters, are you seeing a growing mix of white-box networking switch deployments? I ask because while your custom sales and business continues to broaden, it will be helpful to better understand the growing mix of your $11 billion AI networking portfolio combined this year. Thank you.

Hock Tan

Let me have Charlie address this question. He's the expert.

Charlie Kawwas

Yes. Thank you, Hock. So two quick things on this. One is the – you are exactly right that the portfolio we have is quite unique in providing that flexibility. And by the way, this is exactly why Hock, in his statements earlier on, mentioned that seven out of the top eight hyperscalers use our portfolio. And they use it specifically because it provides that flexibility. So whether you have an architecture that's based on an endpoint and you want to actually build your platform that way or you want that switching to happen in the fabric itself, that's why we have the full end-to-end portfolio. So that actually has been a proven differentiator for us. And then on top of that, we've been working, as you know, to provide a complete network operating system that's open on top of that using SONiC and Psi, which has been deployed in many of the hyperscalers. And so the combination of the portfolio plus the stack really differentiates the solution that we can offer to these hyperscalers. And if they decide to build their own NICs, their own accelerators are custom or use standard products, whether it is from Broadcom or other, that platform, that portfolio of infrastructure switching gives you that full flexibility.

Karl Ackerman

Thank you.

Operator

Thank you. One moment for our next question, and that will come from the line of C.J. Muse with Cantor Fitzgerald. Your line is open.

CJ Muse

Yeah. Good afternoon. Thank you for taking my question. I was hoping to ask two part software question. So excluding VMware, your Brocade, CA, and Symantec business now running $500 million higher for the last two quarters. So curious, is that the new sustainable run rate or were there onetime events in both January and April that we should be considering? And then the second question is as you think about VMware Cloud Foundation adoption, are you seeing any sort of crowding out of spending like other software guys are seeing as they repurpose their budgets to IT? Or is that business so less discretionary that it's just not an impact to you? Thanks so much.

Hock Tan

Well, on the second one, I don't know about any crowding out, to be honest. It's not. What we are offering, obviously, is not something that they would like to use themselves, to be able to do themselves, which is they're already spending on building their own on-prem data centers. And typical approach people take, a lot of enterprises take historically continue today than most people do a lot, people do is they have best of breed. What I mean is they create a data center that is compute as a separate category, best compute there is and they often enough use vSphere for compute virtualization due to improved productivity, but best of breed there. And best of breed on networking and best of breed on storage with a common management operations layer, which very often is also VMware we realize. And what we're trying to say is this mixed bag, and what they see -- is this mixed bag best of [big] (ph) data center, very heterogenous, is not driving that, is not a highly resilient data center. I mean, you have a mixed bag. So it goes down. Where do you find quickly root cause? Everybody is pointing fingers at the other. So you got a problem, not very resilient and not necessary secure between bare metal in one side and software on the other side. So it's a natural thinking on the part of many CIOs we talk to, to say, hey, I want to create one common platform as opposed to just [best-of-breed of age] (ph). So that gets us into that. So it is a greenfield that’s not bad, they started from scratch. If it's a brownfield, that means they have existing data centers trying to upgrade. It's -- that sometimes that's more challenging for us to get that adopted. So I'm not sure there's a crowding out here. There's some competition, obviously, on greenfield, where they can spend their budget on an entire platform versus best-of-breed. But on the existing data center where you're trying to upgrade, that's a trickier thing to do. And it cuts the other way as well for us. So that's how I see it. So in that sense, best answer is I don't think we're seeing a level of crowding out that is -- any and that

very significant for me to mention. In terms of the revenue mix, no, Brocade is having a great, great field year so far and still chugging along. But will that sustain? Hell no, you know that. Brocade goes through cycles like most enterprise purchases. So we're enjoying it while it lasts.

CJ Muse

Thank you.

Hock Tan

Thanks.

Operator

Thank you. And we do have time for one final question, and that will come from the line of William Stein with Truist Securities. Your line is open.

William Stein

Great. Thanks for squeezing me in. Hock, congrats on the yet another great quarter and a strong outlook in AI. I also want to ask about something you mentioned with VMware. In your prepared remarks, you highlighted that you've eliminated a tremendous amount of channel conflict. I'm hoping you can linger on this a little bit and clarify maybe what you did. And specifically also what you did in the heritage Broadcom software business, where I think historically, you've shied away from the channel. And there was an idea that perhaps you'd reintroduce those products to the channel through a more unified approach using VMware's channel partners or resources. So any sort of clarification here, I think, would be helpful.

Hock Tan

Yes, thank you. That's a great question. Yes, VMware taught me a few things. They have 300,000 customers, 300,000. That's pretty amazing. And we look at it. I know under CA, we took a position that let's pick an A-list strategic guy and focus on it. I can't do that in VMware. I approached it differently. And I start to learn the value of a very strong bunch of partners they have, which are a network of distributors and something like 15,000 VARs, value-added resales supported with these distributors. So we have doubled down and invested in this resale network in a big way for VMware. It's a great move, I think but six months into the game. But we are seeing a lot more velocity out of it. Now these resellers, having said that, tend to be very focused on a very long tail of their 300,000 customers. The largest 10,000 customers of VMware are large enterprises who tend to -- they are very large enterprises, the largest banks, the largest health care companies. And their view is I want very bespoke service, support, engineering solutions from us. So we've created a direct approach, supplemented with the VAR of choice where they need to. But on the long tail of 300,000 customers, they get a lot of services from the resellers, value-added resellers, and so in their way. So we now strengthen that whole network of resellers so that they can go direct, manage, supported financially with distributors. And we don't try to challenge those guys unless the customers. On the end of the day, the customer chose where they like to be supported. So we kind of simplify this together with the number of SKUs they have. In the past, unlike what we're trying to do here, everybody is a partner. I mean, you're talking a full range of partners. And whoever makes the biggest deal gets the lowest. The partner that makes the biggest deal gets the lower -- biggest discount, lowest price. And they are out there basically kind of creating a lot of channel chaos and conflict in the marketplace. Here, we don't. The customers, I am aware. They can take it direct from VMware to their direct sales force or they can easily move to the resellers to get it that way. And as a third alternative which we offer, if they chose not -- they want to run their applications on VMware and they want to run it efficiently on a full stack. They have a choice now of going to a hosted environment managed by network of managed service providers, which we set up globally, that will run the infrastructure, invest and operate the infrastructure. And these enterprise customers just run their workloads in and get it as a service, basically VMware as a service. That's a third alternative, and we are clear to make it very distinct and differentiated for our end-use customers. They're available to all three is how they choose to consume our technology.

William Stein

Great. Thank you.

Operator

Thank you. I would now like to hand the call over to Ji Yoo, Head of Investor Relations, for any closing remarks.

Ji Yoo

Thank you, Cherie. Broadcom currently plans to report its earnings for the third quarter of fiscal '24 after close of market on Thursday, September 5, 2024. A public webcast of Broadcom's earnings conference call will follow at 2:00 p.m. Pacific Time. That will conclude our earnings call today. Thank you all for joining. Operator, you may end the call.

Operator

Thank you all for participating. This concludes today's program. You may now disconnect.