AMZN Earnings Call – FY2025 Q2

Generated by discountingcashflows.com

Date: July 31, 2025

Operator

Thank you for standing by. Good day, everyone, and welcome to the Amazon.com Second Quarter 2025 Financial Results Teleconference. [Operator Instructions] Today's call is being recorded. And for opening remarks, I will be turning the call over to the Vice President of Investor Relations, Mr. Dave Fildes. Thank you, sir. Please go ahead.

Dave Fildes

Hello, and welcome to our Q2 2025 financial results conference call. Joining us today to answer your questions is Andy Jassy, our CEO; and Brian Olsavsky, our CFO. As you listen to today's conference call, we encourage you to have our press release in front of you, which includes our financial results as well as metrics and commentary on the quarter. Please note, unless otherwise stated, all comparisons in this call will be against our results for the comparable period of 2024. Our comments and responses to your questions reflect management's views as of today, July 31, 2025, only and will include forward-looking statements. Actual results may differ materially. Additional information about factors that could potentially impact our financial results is included in today's press release and our filings with the SEC, including our most recent annual report on Form 10-K and subsequent filings. During this call, we may discuss certain non-GAAP financial measures. In our press release, slides accompanying this webcast and our filings with the SEC, each of which is posted on our IR website. You will find additional disclosures regarding these non-GAAP measures, including reconciliations of these measures with comparable GAAP measures. Our guidance incorporates the order trends that we've seen to date and what we believe today to be appropriate assumptions. Our results are inherently unpredictable and may be materially affected by many factors, including fluctuations in foreign exchange rates, changes in global economic and geopolitical conditions, tariff and trade policies and customer demand and spending, including the impact of recessionary fears; inflation, interest rates, regional labor market constraints, world events, the rate of growth of the Internet, online commerce cloud services and new and emerging technologies and the various factors detailed in our filings with the SEC. Our guidance assumes, among other things, that we don't conclude any additional business acquisitions, restructurings or legal settlements. It's not possible to accurately predict demand for our goods and services, and therefore, our actual results could differ materially from our guidance. And now I'll turn the call over to Andy.

Andrew R. Jassy

Thanks, Dave. Today, we're reporting $167.7 billion in revenue, up 12% year-over-year, excluding the impact from foreign exchange rates. Operating income was $19.2 billion, up 31% year-over-year and trailing 12-month free cash flow was $18.2 billion. We saw good progress across our various customer experiences and businesses this past quarter. Starting with stores. We feel good about both the inputs and outputs of the business. At Amazon, we think of our business in terms of inputs and outputs. Outputs or metrics like revenue are operating margin. But of course, you can't manage at the output level, it's the inputs that drive the outputs. So we spent virtually all of our time internally talking about and goaling against inputs. The inputs that matter most to customers in our stores business are selection, low prices and speed of delivery. We've taken another step forward in selection these past few months headlined by the much requested return of Nike's products to Amazon's retail store. We've added premium brands like Away, Aveda, Marc Jacobs Fragrances and brands from Saks and Amazon like Dolce&Gabbana, Etro, Stella McCartney, Rosetta Getty and La Prairie. And we started expanding our very successful Perishables Pilot, where we offer customers perishables as a point of purchase when they're ordering other items that will be delivered same day from our same-day fulfillment nodes. We're seeing strong customer adoption as 75% of customers who viewed as service this year are first-time shoppers for perishables on Amazon, with 20% of customers who use the service returning multiple times within their first month. Our prices continue to be low and sharp for customers. It's one of the reasons our everyday Essentials growth outpaced the rest of the business globally and representing one out of every 3 units sold. It's also why well-known research firm Profitero has concluded for 8 years in a row that Amazon has the lowest prices of any U.S. retailer, but perhaps the clearest outputs are the rate at which our stores business grew this past quarter and the success we saw in our recent Prime Day event. This year's Prime Day was our biggest ever with record sales, number of items sold and number of prime sign-ups in the 3 weeks leading up to the Prime Day. Customers save billions of dollars in independent sellers, most of which are small and medium-sized businesses saw their best sales performance of any Prime Day event yet. There continues to be a lot of noise about the impact that tariffs will have on retail prices and consumption. Much of it thus far has been wrong and misreported. As we said before, it's impossible to know what will happen. Where will tariffs finally settle especially China? What happens when we deplete the inventory we forward bought or that our selling partners who were deployed in advance of the tariffs going into effect. If costs end up being higher, who will absorb them, but what we can share is what we've seen thus far, which is that through the first half of the year, we haven't yet seen diminishing demand nor price is meaningfully appreciating. We also have such diversity of sellers in our marketplace over 2 million sellers in total with different strategies of whether to pass on higher cost to consumers, the customers are advantaged shopping at Amazon because they're more likely to find lower prices on the items they care about. Further improving delivery speed remains a key focus, and we continue to make progress. We've previously shared how we rearchitect our U.S. inbound network into a regional structure, allowing us to place inventory and ship from locations closer to customers, improving speed and low run costs. That work is delivering tangible results. In Q2, we increased the share of orders moving through direct lanes where packages go straight from fulfillment to delivery without extra stops by over 40% year-over-year. We've also reduced the average distance packages traveled by 12% and lowered handling touches per unit by nearly 15%. We've made progress on order consolidation with more products positioned locally, we're able to pack more items into each box and send fewer packages per order. That has helped drive higher units per box and improved overall cost to serve. Taken together, these improvements are making the network faster and structurally more efficient. We've also set another global speed record in Q2 delivering to prime members at our fastest speeds ever. In the U.S., we delivered 30% more items same day or next day than during the same period of last year. Items customers used to pick up locally and nearby physical stores are now arriving at their door often within hours. And we're working to further improve delivery speeds no matter where customers live, we've recently announced plans to expand our same-day and next-day delivery to tens of millions of U.S. customers and more than 4,000 smaller cities, towns and rural communities by the end of the year. Today, it's already available in more than 1,000 of these communities across the U.S. The early response from customers in these areas have been very positive. They're shopping more frequently and purchasing household essentials and meaningfully higher rates. Automation and robotics are also important contributors to improving cost efficiencies and driving better customer experiences over time. We deployed our 1 millionth robot across our global fulfillment network and unveiled innovations in our last-mile innovation center, such as automated package sorting and a transformative technology that brings packages directly to employees in an ergonomic height. We rolled out DeepFleet, our AI improves robot travel efficiency by 10%. At our scale, it's a big deal. DeepFleet acts like a traffic management system to coordinate robots' movements to find optimal paths and reduce bottlenecks. For customers, it means faster delivery times and lower costs for our team members, our robots handle more of the physically demanding tasks, making our operations network even safer. This combination of robotics and generative AI is just getting started. And while we've made significant progress, it's still early with respect to what will roll out in the next few years. Moving on to Amazon ads. We're pleased with the strong growth, generating $15.7 billion of revenue in the quarter, growing 22% year-over-year. We continue to see strength across our broad portfolio of full funnel advertising offerings that in the U.S. alone help advertisers reach an average ad-supported audience of more than 300 million across our own properties. These are properties like our retail marketplace, Prime Video, Twitch and Fire TV in live sports, such as NFL, NASCAR and the NBA as well as third-party websites and apps. Another area we're excited about is our demand side platform, or Amazon DSP. Our DSP enables advertisers to plan, activate and measure full funnel investments. Our trillions of proprietary browsing, shopping and streaming signals paired with extensive supply side relationships and our secure clean rooms provide advertisers the ability to optimize advertising, deliver greater precision and drive efficient and effective advertising outcomes. And in June, we announced a momentous partnership with Roku, giving advertisers access to 80 million connected TV households. The largest authenticated connected TV footprint in the U.S. exclusively through Amazon DSP is a giant leap forward for advertisers bringing best-in-class planning, audience precision and performance to TV advertising. We also announced an integration between Disney's real-time ad exchange and Amazon DSP. This collaboration allows advertisers to gain direct access to Disney's premium inventory across platforms like Disney+, ESPN and Hulu while allowing them to leverage insights from both companies. When advertisers work with Amazon, they're not just buying ad space they're benefiting from exceptional programming, innovative technology and unrivaled signals, measurement and audience development that provide strong relevancy for consumers and return on investment for brands. Moving on to AWS. In Q2, AWS grew 17.5% year-over-year and now has over $123 billion annualized revenue run rate. We continue to help organizations of all sizes accelerate their transition to the cloud, signing new agreements with companies, including PepsiCo, Airbnb, Peloton, NASDAQ, London Stock Exchange, Nissan Motor, GitLab, SAP, Warner Bros. Discovery, 12 Labs, FICO, Iberia Airlines, SK Telecom and NatWest. In the rapidly evolving world of generative AI, AWS continues to build a large fast-growing triple-digit year-over-year percentage multibillion-dollar business with more demand than we have supplied for at the moment. A few points to make. First, on the hardware side, our custom AI chip, Trainium2 is landing capacity in larger quantities and has impressively emerged as the backbone for Anthropic's newest generation Claude models and many of our most essential offerings like Amazon Bedrock. We've also launched Amazon EC2 instances powered by NVIDIA Grace Blackwell Superchips, AWS' most powerful NVIDIA GPU accelerated instance. Second, in Bedrock, we've recently added Anthropic's Claude 4 and is the fastest-growing model ever in Bedrock. We've also continued to see strong adoption of Amazon Nova, our own frontier model, and it's now the second most popular foundation model in Bedrock. New features in Nova allow customers to customize their Nova models in ways they can't and other foundation models, allowing organizations to infuse these models with their unique expertise while optimizing for cost and speed. As people have become excited about building agents, they're realizing they lack the tools to build them. In May, we released Strands, an open-source way to more easily build agents has taken off with a wide range of customers with already 2,500 stars on GitHub and over 300,000 downloads on [indiscernible] PI. Customers are also struggling with deploying agents into production in a secure and scalable way. It's holding up enterprises scaling agents. To help solve that problem, Bedrock just released AgentCore. AgentCore is a set of building blocks that gives customers the industry's first secure serverless run time to provide both synchronous and asynchronous execution, aging identity and boundaries, a memory service a gateway to translate services to MCP compatible interfaces, built-in code execution and web browser tools and an observability service. Customers are excited about AgentCore, and it frees them up to start deploying agents more expansively. Third, you're starting to see AWS release more powerful applications at the top layer of the AI stack. AWS transform as an AWS agent that dramatically reduces mainframe modernization time lines from years to months completes VMware TC2 conversions up to 80x faster, it makes it simple to move from .NET windows to .NET Linux implementations, reducing licensing costs for .NET applications by up to 40%. We've also just released Kiro, our new Agentic integrated development environment coding agent. There's a lot of buzz around Kiro with several hundred thousand developers using and requesting access in the first couple of weeks, 100,000 used in the first 5 days of the preview. What struck a cord for developers is that Kiro allows them to do Vibe Coding where developers use natural language to chat with a coding agent to build code. But unlike other coding agents, where developers don't really have any structure to build on top of, Kiro allows developers to use natural language to build spec and then automatically updates that spec as they continue to vibe code or interact with Kiro. This makes it much easier to go from prototyping to production. Customers also like Kiro's event-driven agent hooks that act like an experienced developer catching things developers might miss. When developers save a react component, hooks update that test file. When they modify API endpoints, Hooks refresh readme files. When they're ready to commit security hook scan for leak credentials. It's still very early for Kiro, but it seems clear we're on to something customers love and Kiro has a chance to transform how developers build software. I say this frequently, but remember that 85% to 90% of worldwide IT spend is still on-premises versus in the cloud. In the next 10 to 15 years, that equation is going to flip, further accelerated by company's excitement for leveraging AI. So AWS has significantly broader functionality, stronger security and operational performance a much deeper experience helping enterprises modernize their infrastructure bodes well for the AWS business moving forward. We're also seeing momentum in a number of our other areas across Amazon. I'll mention just a few. We're excited about our progress with Alexa+, our next-generation assistant powered by generative AI. We've been rolling out early access to U.S. customers to start millions of customers have access now. We're seeing very positive feedback, and we'll continue to iterate on the experience. We've recently completed our third successful launch of Project Kuiper. We haven't launched this service commercially yet but already have an impressive amount of enterprise and government customers who have signed agreements to use Kuiper. In Prime video live sports, our first season of NASCAR drew about 2 million viewers per race and the youngest audience among NASCAR broadcasters in more than a decade. We've recently announced our stellar broadcasting crew for our upcoming first NBA season, including in Eagle, Stan Van Gundy, Kevin Harlan, Dwayne Wade, Taylor Rooks, Blake Griffin, Dirk Nowitzki, Steve Nash and Candace Parker. We also announced Denis Villeneuve, an Academy Award nominee, as the Director for the next James Bond film, James Bond is in the hands of one of today's greatest filmmakers, and we cannot wait to get started on 007's next adventure. Finally, we continue to be very pleased with the growth and residence of Amazon Pharmacy and it's grown 50% year-over-year, year- to-date on an already significant size base. A lot of good things happening across the company. With that, I'll turn it over to Brian for a financial update.

Brian T. Olsavsky

Thanks, Andy. Let's start with our top line financial results. Worldwide revenue was $167.7 billion, a 12% increase year-over-year, excluding the impact of foreign exchange. Foreign exchange had a $1.5 billion favorable impact to revenue in the quarter as foreign currencies generally strengthened versus the U.S. dollar. As a reminder, our Q2 revenue guidance had anticipated an unfavorable impact of approximately 10 basis points or $100 million. Worldwide operating income was $19.2 billion, which was $1.7 billion above the high end of our guidance range. Across our segments, we continue to prioritize cost-effective innovation that delivers value for our customers. In the North America segment, second quarter revenue was $100.1 billion, an increase of 11% year-over-year. International segment revenue was $36.8 billion, an increase of 11% year-over-year excluding the impact of foreign exchange. Worldwide paid units grew 12% year-over-year. We remain focused on inputs that matter most to our customers. In the second quarter, we saw broad-based strength across our key performance metrics. This includes sharp pricing and more in stock availability as well as record delivery speeds for Prime members. Our millions of global sellers continue to be an important contributor to our vast selection. This helps customers find the items they need and does so at a competitive price. Our investment in tools, services and fast delivery speeds help our selling partners reach more customers and further scale their businesses. In Q2, worldwide third-party seller unit mix was 62%, the highest ever, up on 100 basis points from Q2 of last year. We're also closely monitoring the macroeconomic environment, including the impact of tariffs. As Andy mentioned, our Q2 plan factored in a range of assumptions, not all of which materialized. We will continue to consider a range of assumptions going forward. Shifting to profitability. North America segment operating income was $7.5 billion, an increase of $2.5 billion year-on-year. North America operating margin was 7.5%, up 190 basis points year-over-year. International segment operating income was $1.5 billion, up $1.2 billion year-over-year. International operating margin was 4.1%, up 320 basis points year-over-year. We're pleased with the strong execution of our operations teams and the positive experience they delivered for customers. In Q2, we saw productivity gains in our transportation network, driven by improved inventory placement, strong leverage on high unit volumes and higher levels of in-demand inventory from both first-party and third-party selling partners. These factors contributed to faster delivery speeds and lower costs. Outbound shipping costs were up 6% year-over-year and continue to grow at a meaningfully slower pace than unit growth, which as I mentioned earlier, was up 12% year-over-year. We're committed to initiatives that further improve our cost structure. Strategic inventory placement drives multiple benefits, including better in-stock availability, shorter delivery routes and faster customer delivery times. When we optimize inventory location, we can consolidate more items for package, reducing packaging materials and costs. To achieve this, we will continue to improve upon our inbound network, expand our U.S. same-day delivery facilities, including rural communities and implement robotics and automation across our facilities. While year-over-year improvements in operating margin may fluctuate, we have a purposeful strategy to achieve sustained progress over time. Shifting to advertising. Advertising revenue grew 22% year-over-year, driven by sponsored products as we saw strong traffic in our stores. Advertising remains an important contributor to profitability in the North American International segments. Our full funnel advertising approach of connecting brands with customers is resonating. Moving next to our AWS segment. Revenue was $30.9 billion, an increase of 17.5% year-over-year. AWS now has an annualized revenue run rate of more than $123 billion. During the second quarter, we continue to see growth in both our generative AI and non- generative AI businesses as companies turn their attention to newer initiatives bring more workloads to the cloud, restart or accelerate existing migrations from on-premise to the cloud and tap into the power of generative AI. AWS operating income was $10.2 billion. We did see AWS segment margins decline from a record high of 39.5% in Q1 to 32.9% in Q2. The largest quarter-over-quarter driver of the decrease or about half is due to the seasonal step-up in stock-based compensation expense, driven by the timing of our annual compensation cycle. AWS margins also saw headwinds from higher depreciation expense as well as unfavorable impacts from year-over-year fluctuations in foreign exchange rates. The depreciation expense is a result of our growing investments in capital expenditures in AWS. As we've said in the past, we expect AWS operating margins to fluctuate over time driven in part by the level of investments we are making at any point in time. We will continue to invest more capital in chips, data centers and power to pursue this unusually large opportunity that we have in generative AI. Now turning to our cash CapEx, which was $31.4 billion in Q2. We expect Q2 CapEx to be reasonably representative of our quarterly capital investment rate for the back half of this year. AWS continues to be the primary driver as we invest to support demand for our AI services and increasingly in custom silicon, like Trainium, as well as tech infrastructure to support our North America and international segments. Additionally, we continue to invest in our fulfillment and transportation network to support growth of the business, improve delivery speeds and lower our cost to serve by investing in same-day delivery facilities as well as robotics and automation. Collectively, these investments will support growth for many years to come. Moving on to our third quarter financial guidance. As a reminder, our guidance considers a range of possibilities, which take into consideration Q2 results, trends we see quarter-to-date and expectations around the macroeconomic environment, including tariffs. Q3 net sales are expected to be between $174 billion and $179.5 billion. We estimate the year-over-year impact of changes in foreign exchange rates based on current rates, which we expect to be a favorable impact of approximately 130 basis points. As a reminder, global currencies can fluctuate during the quarter. Q3 operating income is expected to be between $15.5 billion and $20.5 billion. In this dynamic environment, we'll focus on what matters most, delivering exceptional customer value through product selection, competitive prices and unmatched convenience. We remain focused on driving a better customer experience and believe putting customers first is the only reliable way to create lasting value for our shareholders. With that, let's move on to your questions.

Operator

[Operator Instructions] Our first question comes from Doug Anmuth with JPMorgan.

Douglas Till Anmuth

I have two. First, can you just help us understand with some more granularity how tariffs are being absorbed across suppliers, Amazon and consumers and whether you anticipate any change going forward? And then second on AWS, we're seeing significantly faster cloud growth among the #2 and #3 players in the space. I totally appreciate that AWS is coming off of a bigger base. But beyond that, do you think the output gap is due more to customer demand or infrastructure supply for both?

Andrew R. Jassy

Yes. I'll take both of those. I'll start with the tariffs. I think what we said a number of times, and we still believe it is we just don't know what's going to happen moving forward. It's hard to know where the tariffs are going to settle, particularly in China. It's hard to know what will happen when we deplete some of the prebuys that we did on our own first-party retail and then some of the forward deploying that we saw of our third-party selling partners. And if costs go up over time, it's -- we're unsure at this point who's going to end up absorbing those higher costs. What we can tell you is what we've seen so far in the first half of the year, in the first half, we just haven't seen diminished demand. And we haven't seen any kind of broad scale ASP increases. And so that could change in the second half. There are a lot of things that we don't know, but that's what we've seen so far. On the question on AWS, the first thing I'd say is it's -- as you said, Doug, in your question, you -- year-over-year percentages and growth rates are always a function of the base in which you operate. And we have a meaningfully larger business in the AWS segment than others. I think the second player is about 65% of the size of AWS. And we -- when we look at the results over the last number of quarters, there are sometimes where -- as far as we can tell, we're growing faster than others and sometimes others are growing faster than us. But it's still like if you look at second place player you're talking about, it's a -- it's still a pretty significant segment market segment leadership position that we have. And regardless, these are all really just moments in time. The last week is a moment in time too where -- the reality of what really matters is what customers' experiences are in operating on these platforms. And if you look at what matters to customers, what they care a lot about what the operational performances what the availability is, what the durability is, what the latency and throughput is of the various services. And I think we have a pretty significant advantage in that area. They care a lot about security. If you have data that matters. And for most companies, they're putting data that they really care about in the cloud. The security and the privacy of that data matters a lot, and there are very different results in security in AWS than you'll see in other players. And yes, you could just look at what's happened in the last couple of months, you can just see kind of adventures at some of these players almost every month. And so a very big difference, I think, in security. And then I think a really significant difference in functionality where not just in the core infrastructure, do we have a lot more functionality in our services. But I think if you look at our end-to-end offering in AWS in AI from the bottom of the stack all the way to the top, it's pretty different. So I feel good about the inputs and the services that we're offering to customers across AI as well as non AI and we have more demand than we have capacity right now. So we could be doing more revenue and helping customers more, and we're working very hard on changing that outcome and how much capacity we have, but it's still like if you look at the business, it's $123 billion annual revenue run rate business and it's still early. I mean how often do you have an opportunity that's $123 billion of annual revenue run rate where you say it's still early. It's a very unusual opportunity that we're very bullish about.

Operator

Our next question comes from Mark Mahaney with Evercore.

Mark Stephen F. Mahaney

Okay. I'll stick with AWS to start with. Could you just disclose the backlog number? And then in the past, I know you've talked about the supply constraints and hoping that they will sort of resolve themselves by the back half of the year. Is that still your intention? Anything that suggests that the supply constraints are going to get resolved earlier or later? And then a long-term question on Alexa+ and I've been experimenting with it for a while. Just Andy, when you think about the potential that, that has in terms of increasing engagement, maybe tapping into some services revenue, advertising, maybe a little bit more retail sales per household, you're just reducing friction. Just talk about what -- from a financial perspective, how you think that could play out, how we would maybe see that in the numbers.

Dave Fildes

Mark, this is Dave. I'll just start off to give you the backlog figures. So at the end of the quarter, at June 30, that was $195 billion, so that's up about 25% year-over-year.

Andrew R. Jassy

On the supply constraints as it relates to AWS and what we see there, as I mentioned, we have more demand than we have capacity at this point. And I think that -- and you see some of the constraints and they kind of exist in multiple places, the single biggest constraint is power. But you also see constraints off and on with chips and then some of the components that -- once you have the chips to actually make the servers, the sometimes you have new generations of chips that are a little bit later than they're supposed to be and sometimes you get the chips and they yield you get in making servers isn't what you expect when you get to ramp. So there are a bunch of those pieces today that we're working on. It's really true across the industry today. I don't believe that we will have fully resolved the amount of capacity we need for the amount of demand that we have in a couple of quarters. I think it will take several quarters. But I do expect that it's going to get better each quarter, and I'm optimistic about that. I think on the Alexa question, I think I'd start by saying the Alexa+ experience is so much better than I think our prior Alexa experience. She's much more intelligent than her prior self. She's much more capable and I would say unlike the other chat bots that are out there today who are good at answering questions, but really can't take any action for you. Alexa+ can take a lot of action for you, which is very compelling. So I can ask Alexa to play music for me or play video for me to move my music from one device to another or if I'm listening to a song, that's on -- that's in a movie, I can ask Alexa+ to actually put that movie scene on -- of the song I'm playing, and it will put it on my Prime video on Fire TV or if I have guests coming over. I can say, Alexa draw the curtains, put the light on the porch and the driveway, increase the temperature by 5 degrees and put on music that would be great for a dinner party. And she does all that just through using natural language. So she could take a lot of actions and it's compelling. And what we see so far, we've been rolling out Alexa+ starting in the U.S. It's with millions of customers now. The rest in the U.S. coming in the next couple of months and it's starting the international rollout more broadly later in the year. And customers really like the experience. They recognize how much better it is than what it was before. The ratings are very high. The usage is much more expansive than what they were using before and the number of calls they're making is meaningfully higher. And I think there are a number of different areas where we'll see benefit. I think first, if you build the world's best personal assistant, that has a lot of utility for customers, and therefore, it gets used a lot. So it means everything from people are excited about the devices that they can buy from us that has Alexa+ enabled in it. People do a lot of shopping and it's really -- it's a delightful shopping experience that will keep getting better. I think over time, there will be opportunities as people are engaging more multiturn conversations to have advertising play a role to help people find discovery and also as a lever to drive revenue. And I think over time, you could also imagine, as we keep adding functionality that there could be some sort of subscription element beyond what there is today. Today, Prime members get Alexa+ for free and non-Prime members pay $9.99 a month for Alexa+. So I think it's very -- it's still very early days, but we're very encouraged by the experience we're providing and you can bet we're going to be iterating on it constantly.

Operator

And our next question comes from the line of Colin Sebastian with Baird.

Colin Alan Sebastian

I guess, first off, with respect to the International segment and the progress in both revenues and margins. I was hoping you could add maybe some color on the drivers of each of those and the sustainability of the improved efficiency in those markets. And then, Andy, you mentioned Kuiper'. We haven't heard as much about that of late. So maybe if you could review where things stand relative to next year's launch target timing of the service rollout and maybe any perspective you have on the longer-term ambitions for the satellite network.

Brian T. Olsavsky

Colin, this is Brian. I'll start with the international segment question. So a very strong quarter for International segment, both on revenue growth and also on operating margin. Operating margin was up 320 basis points year-over-year to 4.1%. And if you look, that's the continuation of a strong progression we've seen quarter-by-quarter over the last 10 quarters in total, we've seen an increase of close to 700 basis points during that time period. So really, it's a tale of 2 pieces of segment -- or excuse me, sections within international. One is the the established countries like U.K., Germany and Japan, already at similar margin profiles to the U.S. So we'll -- as they continue to grow their contribution of profit will grow over time, and that's what we're seeing. In the quarter, we saw strong productivity in our transportation network in those countries, much like we saw in the U.S., and that's led to higher units for package and faster delivery speeds at lower cost. So again, the theme of lower cost to serve and also increased speed and better selection for our customers continues to grow internationally as well. In our emerging countries, we're pleased with the progress we're making there. We've launched 8 countries, of course, in the last 5 years. And there are all varying degrees of upfront investment in different point in times in their journey to profitability, but they're all making very nice improvements quarter- over-quarter in growing selection, adding prime members and expanding our customer offerings. So I think what you're saying, again, is sustained improvement in those areas. I feel very good about it. The established countries are continuing to grow and develop and very much look like the United States and the emerging countries, again, are all very different stages of growth right now.

Andrew R. Jassy

On the Project Kuiper question, so Project Kuiper's or low Earth orbit satellite constellation that we're putting up and launching and there's 400 million to 500 million households worldwide who don't have broadband connectivity. And it means they can't do a lot of things we take for granted, like education online or business online or shopping or entertainment. There is really a digital divide, and it's much needed. And it's also true for enterprises and for governments that they have assets or needs to have visibility or connectivity that they can't get today given the lack of broadband and a bunch of places around the world. So there's a high need. I would say that as we get our constellation into space, there will really be 2 players that have what I would consider the modern technology in low Earth orbit satellite. One is the incumbent in the market today and the second will be Project Kuiper. I think that we will have a pretty meaningful differentiation here in performance. If you look at the performance of what I expect on the uplink and downlink, I think Project Kuiper will be advantaged. I also think the pricing is going to be very compelling for customers. And then I think if you think about the 3 key customer segments who want low Earth orbit satellite, consumers enterprises and governments. We have very strong relationships with all 3 customer segments given our consumer businesses and our AWS business, and I think if you think about enterprises and governments, a lot of what they want to do when they take the data down from space as they actually want to put it into a cloud to do analysis, analytics and AI and various operations on top of it. And the fact that Project Kuiper and AWS are so seamlessly connected is very attractive to enterprises and to governments. And I'm kind of amazed we haven't launched Project Kuiper yet. But the number of enterprise and government agreements that have been signed already to use Project Kuiper is impressive. So we're working very hard to get the satellites into space. We have some delays with some of the rocket providers, but we have most of the available rocket launches over the next couple of years and were very helpful to get the service into commercial -- into commercial data later this year or early next year.

Operator

And our next question comes from the line of Brian Nowak with Morgan Stanley.

Brian Thomas Nowak

Andy, I have two for you on AWS. They're a little tough, but I'm going to throw at you. So there is a Wall Street finance person narrative right now, that AWS is falling behind in GenAI with concerns about share loss to peers, et cetera. Can you just sort of address that? What is your rebuttal to that? And talk to us about your and the team's most important focal points. Just to ensure that AWS stays on the nice edge of innovation versus hyperscaler peers? And then secondly, I know AWS is a big business. But is there any reason to assume it shouldn't accelerate in the back half and into '26 given the size of the opportunity and all of the GenAI capabilities that are going to sort of come to us in the next 12 months?

Andrew R. Jassy

Yes. So on the first one around AI. The first thing I would say is that I think it is so early right now in AI. If you look at what's really happening in the space, you have -- it's very top heavy. So you have a small number of very large frontier models that are being trained that spend a lot on computing, a couple of which are being trained on top of AWS and others are being trained elsewhere. And then you also have, I would say, a relatively small number of very large-scale generative AI applications. The one category would be chatbots with the largest by a fair bit being ChatGPT, but the other category being really, I'll call it, coding agents. So these are companies like Cursor, Vercel, Lovable and some of the companies like that. Again, several of which run significant chunks on top of AWS. And then you've got a very large number of generative AI applications that are in pilot mode -- or they're in pilots or that are being developed as we speak and a very substantial number of agents that also people are starting to try to build and figure out how to get into production in a broad way, but they're all -- they're quite early. And many of them that are out there are they're significant, but they're just smaller in terms of usage relative to some of those top heavy applications I mentioned earlier. We have a very significant number of enterprises and startups who are running applications on top of AWS' AI services and then -- but they're all -- again, like the amount of usage and the expansiveness of the use cases and how much people are putting them into production and the number of agents that are going to exist. It's still just earlier stage than it's going to be and so then when you think about what's going to matter in AI, what's going to -- what are customers going to care about when they're thinking about what infrastructure use, I think you kind of have to look at the different layers of the stack. And I think for those that are -- both building models, but also just -- if you look at where the real costs are, they're going to ultimately be an inference today, so much of the cost in training because customers are really training their models and trying to figure out to get the applications into production. But at scale, 80% to 90% of the cost will be an inference because you only train periodically, but you're spinning out predictions and inferences all the time. And so what they're going to care a lot about is they're going to care about the compute and the hardware they're using. And we have a very deep partnership with NVIDIA and will for as long as I can foresee, but we saw this movie in the CPU space with Intel, where customers are anchoring for better price performance. And so we built just like in the CPU space, where we built our own custom silicon and building Graviton which is about 40% more price performance than the other leading x86 processors. We've done the same thing on the custom silicon side in AI with Trainium and our second version of Trainium2 is really -- it's become the backbone of Anthropic's next Claude models they're training on top of, and it's become the backbone of Bedrock and the inference that we do. So I think a lot of the inference, it's about 30% and 40% better price performance than the other GPU providers out there right now, and we're already working on our third version of Trainium as well. So I think a lot of the compute and the inference is going to ultimately be run on top of Trainium2. And I think that, that price performance is going to matter to people as they get to scale. Then I would say that middle layer of the stack are really -- it's a combination of services that customers care about to be able to build models and then to be able to leverage existing leading frontier models and then build high-quality generative AI applications that do inference at scale. And we see it for people building models, they continue to use SageMaker AI very expansively, and then Bedrock, when you're leveraging leading frontier models is also growing very substantially. And as I said in my opening comments, the number of agents of scale is still really small in the scheme of what's going to be the case, but part of the problem is it's actually hard to actually build agents. And it's hard to deploy these agents in a secure and scalable way. And so I think the launches we made recently in strands that make it much easier to build agents and then AgentCore that make it much easier to deploy at scale and in a secure way are being very well received and customers are excited is going to change what's possible on the agent side. And then I think that it's -- you've got a very large number. I mean, remember, 85% to 90% of the global IT spend is still on-premises. If you believe that equation is going to flip, which I do, and we do. You have a lot of legacy infrastructure that you've got to move. These are mainframes. These are VMware's instances and when we build agents like AWS Transform to make it much easier to move mainframe to the cloud, much easier to move VMware to the cloud, much easier to move .NET windows to .NET Linux to save money, those are compelling for enterprises or things like Kiro that allow customers to develop in a much easier way and in a much more structured way, which is why I think people are excited about it. So I think I really like the inputs and the set of services that we're building in the AI space today. Customers really like them and it's resonating with them. I still think it's very early days in AI and in terms of adoption. But the other thing I would just say is that. Remember, because we're at a stage right now where so much of the activity is training and figuring out how to get your generative AI applications into production. People aren't paying as close attention as they will and making sure that those generative AI applications are operating where the rest of their data and infrastructure. Remember, a lot of generative AI inference is just going to be another building block like compute, storage and database. And so people are going to actually want to run those applications close to where the other applications are running, where their data is. There's just so many more applications and data running in AWS than anywhere else. And I'm very optimistic about as we get to a bigger scale what's going to happen to AWS on the AI side. And I think we have a set of services that is unique top to bottom in the stack. I think on the last part about what do we expect with respect to acceleration, we don't give guidance by segment. So I'm not going to try and give you guidance. But I do believe that the combination of more enterprises who have resumed their march to modernize their infrastructure and move from on-premises to the cloud, coupled with the fact that AI is going to accelerate in terms of more companies deploying more AI applications into production that start to scale, coupled with the fact that I do think that more capacity is going to come online in the coming months and quarters, make me optimistic about the AWS business.

Operator

And the next question comes from the line of Ron Josey with Citi.

Ronald Victor Josey

That was really helpful. Maybe taking that same question, but focusing internally on Amazon, I think you penned an article or a blog post back in June, just talking about the ability or potential with GenAI agents and improving efficiencies and speed to market internally. So would love to hear your thoughts as you think about how Amazon is adopting generative AI internally, how perhaps you've seen improving speed to market as a result of everything you've just talked about.

Andrew R. Jassy

Yes. I think that AI is the biggest technology transformation for a lifetime, which is saying a lot because even in some of our relatively short lifetimes, we've been through the cloud, mobile and the Internet itself. But I think it's going to be the biggest transformation technically in our lifetime. And I think it's going to not only change every customer experience that we know and enable customer experience we really only dreamed about before, but it's also going to change very substantially the way we work. And if you think about it, the way that we do coding, the way that we do analytics, the way that we do research, the way that we do finance and measure -- I mean, really, the way we do business process automation, the way we do customer service. Every single area that I can think of in the way we work is likely going to be impacted in some meaningful way by AI. And I think when you have a big shift like that, you have 2 macro choices. You can either decide that you're going to embrace it. And you're going to help shape it and you're going to figure out how to build the right tools to allow you to take advantage of the technology or you can wish it away and have it shape you. And the posting that you're referencing, Ron, that I made was just really being clear with the team that we're going to pursue that former approach. We are going to embrace it. We're going to try and shape it. So we have a number of tools and agents that we've built already inside the company. These are things like if you think about Kiro and the opportunity to have coding agents do a lot of a coding that's very compelling. It's going to allow our teammates to be able to start from a more advanced starting spot and to be able to invent for customers much more quickly and much more expansively. If you think about services like Connect, which is our AWS service, it does call center software that has a lot of AI built into it that changes the productivity of all your customer service agents. And you can just imagine across the board, the types of things we're going to do to make it easier to get software out to build high- quality software to do operations to automate a lot of the business process coordination that happens inside the kind. We're going to work on a whole bunch of those areas to allow ourselves primarily to be able to invent for customers much more quickly and expansively. But also, I think it's going to make all of our teammates jobs much more enjoyable because they won't have to do as many of the road parts of the job that we all do right now, and they're going to be able to own more pieces of what they're trying to solve for customers. And we want deep owners that want to own as much end-to-end as possible.

Operator

And our final question will come from the line of Justin Post with Bank of America.

Justin Post

Great. I'll just ask on the revenue guidance. It looks pretty robust growth in the third quarter. I know you won't say whether AWS is expected to accelerate. But could you talk about some of the drivers and what kind of tariff and other contingencies you've put in there? And then maybe any thoughts on how you're thinking about how the Q4 is shaping up.

Brian T. Olsavsky

Justin, this is Brian. I'll take this one. We've guided to $174 billion to $175 billion. And we -- excuse me, $179.5 billion is a typo. We feel good about the growth rate in Q2 that we had and the acceleration in a number of areas, including units. And we had a very successful Prime Day earlier this month. So there is uncertainty on where tariffs will settle and maybe the ultimate impact on consumers, as Andy mentioned earlier. But we feel really good about the key inputs we control price selection and convenience. We've talked about delivery speeds increasing. We've talked about selection increasing high-end stock levels, well dispersed inventory close to customers. So we think that all works in our favor. So I would say we're cautiously optimistic I'm not going to give any guidance for Q4 at this time, but we'll talk about that next time.

Dave Fildes

Thanks for joining us on the call today and for your questions. A replay will be available on our Investor Relations website for at least 3 months. We appreciate your interest in Amazon, and we look forward to speaking with you again with you next quarter.

Operator

And ladies and gentlemen, that does conclude today's teleconference. You may disconnect your lines at this time. Thank you for your participation.